

When Joint Training with Machine-Generated Text Detection Can Distort Sarcasm Detection:

A Bilingual Diagnosis of Class-Conditional Negative Transfer

Anonymous ACL submission

Abstract

Sarcasm is non-literal language whose intended meaning may diverge from apparent sentiment, complicating detection in online text. As LLM-generated content appears alongside human writing, systems may need to judge both sarcasm and machine authorship. These judgments interact in generation, where detector scores may be treated as measures of sarcastic intent and human authorship. Their interaction in shared representations remains insufficiently understood. MGTD may promote source-predictive stylistic or distributional cues, whereas sarcasm detection should reflect pragmatic intent rather than source. We conduct, to our knowledge, the first bilingual diagnosis of how joint training with MGTD changes sarcasm detection across Chinese and English. We harmonize bilingual resources, generate topic-matched AI texts, and build DualTest with source and sarcasm labels, defining AI-side sarcasm labels by requested generation style. Same-architecture joint and single-task procedures reveal class-conditional negative transfer concentrated on AI-generated text prompted to be non-sarcastic. This pattern persists with a second generator, while substantially lower rates are observed in the evaluated RCNN and interaction-projection configurations. An exploratory comparison shows that detector scores for human authorship and LLM ratings of perceived human-likeness rank generated outputs differently. These findings motivate subgroup evaluation and caution against detector-only assessment.

1 Introduction

A sarcastic sentence typically conveys an attitude or intent that differs from, and may oppose, its literal meaning. Sarcasm can invert sentiment polarity and disrupt opinion-mining systems (Tay et al., 2018); ironic or sarcastic abuse is also a known challenge for content moderation (Vidgen et al., 2019). The task remains difficult because it de-

pends on contrast, intent, and context. Annotation outcomes vary with elicitation method and annotator perspective (Oprea and Magdy, 2020; Casola et al., 2024), while model behavior is sensitive to domain and style (Jang and Frassinelli, 2024); multilingual settings introduce further variation (Casola et al., 2024) (§2.1).

LLM-generated text adds a source dimension to this task. Human-authored and AI-generated texts can each be sarcastic or non-sarcastic, so all four combinations are possible. A robust sarcasm detector should not infer sarcastic intent from source alone.

Machine-generated text detection (MGTD) imposes a contrasting representational demand: its objective distinguishes human-written from AI-generated text and therefore favors source-predictive representations (Guo et al., 2023; Wang et al., 2024) (§2.2). Some latent stylistic or distributional cues predictive of AI-generated text may overlap with, or become entangled with, features used for sarcasm detection. The output labels do not conflict; the potential conflict lies in how the two tasks organize and use the shared representation. Sarcasm detection should not rely on source differences alone, whereas MGTD must encode them. Shared-encoder multi-task learning can suffer task conflict and negative transfer (Mueller et al., 2022); more specifically, it can exploit knowledge useful for one task but spurious for another (Hu et al., 2022b). Evaluating the tasks only in separate models cannot reveal how incorporating MGTD into a joint-training procedure changes the sarcasm decision.

We study joint training through a controlled same-architecture comparison rather than assuming that an auxiliary task must improve accuracy. A dual-head model lets us test how sarcasm decisions change when MGTD batches additionally optimize the shared representation. Single-head controls provide a direct procedural counterpoint,

and source $\times$ sarcasm subgroups show where the change appears.

As a downstream implication, we examine whether the same detectors provide reliable evaluation signals for generated sarcasm. ViSP, for example, derives PPO rewards from a sarcasm detector (Wang et al., 2026), while generated sarcasm may be judged for both sarcastic intent and human-likeness (§2.4). If detector scores are source-sensitive, they could favor source-predictive stylistic or distributional cues rather than measure sarcastic intent. We therefore compare detector scores for human authorship with source-blinded LLM ratings of perceived human-likeness separately in §6.

We investigate the interaction through same-architecture comparisons. We train dual-head models on bilingual Chinese–English data and build DualTest, whose dual labels partition evaluation into four source $\times$ sarcasm quadrants. Compared with single-head controls using the same sarcasm data, the evaluated joint-training procedure substantially increases the sarcasm false-positive rate on AI-generated text prompted to be non-sarcastic. We refer to this procedure-induced degradation, concentrated in a particular source $\times$ sarcasm subgroup, as class-conditional negative transfer. The pattern persists with a second generator, with substantially lower rates in the evaluated RCNN and interaction-projection configurations than in the plain configuration.

We then examine the shared representation. Source information is highly decodable and correlates with sarcasm decisions. Representation interventions show that the subgroup gap is sensitive to a dominant-variance subspace, while leaving the causal content of that subspace unresolved (§5.3.2, Appendix E). Because probe accuracy alone does not establish how encoded information is used (Hewitt and Liang, 2019; Elazar et al., 2021), we treat the source-correlated shortcut as a mechanism hypothesis with preliminary support rather than an established cause.

Contributions. To our knowledge, we provide the first bilingual diagnostic study of how joint training with MGTD changes sarcasm detection across Chinese and English, together with DualTest for source $\times$ sarcasm analysis. Same-architecture comparisons across models and generators identify class-conditional negative transfer concentrated on AI-generated text prompted to

be non-sarcastic; cross-architecture comparisons characterize how its severity varies across the evaluated configurations. Representation analyses characterize and constrain the source-shortcut hypothesis. Finally, an exploratory detector–judge comparison illustrates a downstream risk of using detector scores to evaluate generated sarcasm and perceived human-likeness.

2 Related Work

2.1 Sarcasm Detection and Multilingual Resources

Feature-engineered methods (Riloff et al., 2013), neural sequence models (Ghosh and Veale, 2016; Tay et al., 2018), and pretrained language models (Khatri and P, 2020; Du et al., 2022) have each advanced sarcasm detection. The task remains difficult in multilingual and class-imbalanced settings, especially because datasets differ in whether labels reflect intended sarcasm or annotator-perceived sarcasm (Oprea and Magdy, 2020; Abu Farha et al., 2022; Casola et al., 2024). Although scores are not directly comparable across datasets, languages, modalities, or evaluation setups, they illustrate substantial benchmark variability. On English iSarcasmEval (Abu Farha et al., 2022), the top two systems report positive-class F1 of 0.605 and 0.569; across MultiPICO languages (Casola et al., 2024), zero-shot ChatGPT positive-class F1 ranges from 0.152 to 0.567. On the bilingual multimodal SarcNet benchmark (Yue et al., 2024), a text-only XLM-R baseline reaches 0.734 F1.

Architecturally, this work follows Yang and Ikeda (2024), who compared mBERT, mBERT-BiLSTM, and mBERT-RCNN on mixed English–Arabic and English–Chinese training sets and found that the RCNN variant generally improved performance on two English test sets. Their RCNN implementation builds on the recurrent convolutional architecture of Lai et al. (2015); we therefore include RCNN among the intermediate architectures examined in §4.1.

The benchmark results summarized above come from different datasets, so §5.1 provides directly comparable mBERT dual-head baselines on DualTest. To our knowledge, existing sarcasm resources do not provide both sarcasm and human-versus-AI source labels. DualTest supplies both labels in a bilingual Chinese–English setting.

The Chinese and English source resources differ in domain, annotation procedure, and class bal-

ance, so bilinguality introduces additional distributional variation. We treat language as a robustness and stratification dimension rather than as the primary transfer objective: the study asks whether the subgroup failure is confined to one language, not whether knowledge transfers across languages. We therefore report language-stratified results in addition to pooled metrics.

2.2 Machine-Generated Text Detection and Authorship Verification

Machine-generated text detection (MGTD) includes zero-shot model-likelihood methods, including probability-curvature detectors (Mitchell et al., 2023), and discriminative models trained on paired human–AI corpora (HC3; Guo et al., 2023). Generalization is a central difficulty: M4 (Wang et al., 2024) provides a benchmark across generators, domains, and languages and shows that detectors struggle on unseen domains or LLMs, while RAID (Dugan et al., 2024) stress-tests detectors under unseen generators, decoding variations, and adversarial attacks. This generalization gap is consistent with reliance on source-predictive regularities that are partly generator- or domain-specific; such cues are relevant here because they may enter a representation shared with sarcasm detection. Bevendorff et al. (2025) distinguish closed-set authorship attribution from open-set authorship verification and argue that realistic LLM detection is better cast as the latter. Our binary MGTD setup is narrower: we use it as a controlled source-classification auxiliary task, not as a measure of text quality.

In our literature search, we found no work that jointly trains sarcasm detection and MGTD with a shared encoder.¹ This gap motivates our study because source information has different roles in the two tasks: it is the target of MGTD but should not by itself determine a sarcasm detector’s prediction. Our study tests whether a shared representation preserves that distinction.

2.3 Multi-Task Negative Transfer, Shortcut Learning, and Shared-Encoder Mitigation

In multi-task learning, auxiliary tasks are not always beneficial: transfer depends on task relat-

edness and alignment between task data (Zhang et al., 2023; Wu et al., 2020), while gradient conflict is a central optimization challenge that can contribute to negative transfer (Yu et al., 2020; Liu et al., 2021). Mueller et al. (2022) provide a relevant NLP precedent through matched single- and multi-task comparisons. They find similar negative-transfer patterns in canonical and text-to-text systems and note that models may infer task identity from input distributions, a strategy that may not transfer reliably out of distribution. He and Choi (2021) and Li et al. (2023) further analyze functional specialization in multi-task Transformer encoders. Shortcut learning offers a complementary reading: models may exploit features that are predictive in the training distribution yet fail under distribution shift (Geirhos et al., 2020; Ye et al., 2026). In MTL, shared representations can increase reliance on spurious correlations (Hu et al., 2022b; Chai et al., 2025). Related multi-dataset work shows that adding data can introduce spurious correlations and reduce worst-group performance (Compton et al., 2023).

Mitigation under a shared encoder takes two broad forms. Parameter-efficient methods introduce task-specific or task-conditioned lightweight components (AdapterFusion, Pfeiffer et al., 2021; MAD-X, Pfeiffer et al., 2020; HyperFormer++, Karimi Mahabadi et al., 2021; PALs, Stickland and Murray, 2019). Other methods explicitly gate or adapt information sharing within the shared computation (Wu et al., 2019; CA-MTL, Pilault et al., 2021).

Rather than isolating task-specific parameters, we examine whether alternative intermediate representations alter the subgroup failure induced by joint training. We compare a no-module baseline with two intermediate-module variants (§5.3). Unlike aggregate negative transfer, our analysis asks whether the degradation under the evaluated joint-training procedure concentrates in particular source $\times$ sarcasm subgroups.

2.4 Misalignment in Detector-Based Evaluation and Rewards

The representational conflict studied above may also matter when detector outputs are reused beyond classification, particularly as evaluation or reward signals for generated text. Automated scorers, including LLM judges, are increasingly used to evaluate generated text (Zheng et al., 2023). Model-based scores are also optimized as rewards,

¹Searches conducted in May 2026 and updated in July 2026 across Semantic Scholar, arXiv, ACL Anthology, and Google Scholar, using keyword combinations including “sarcasm detection,” “machine-generated text detection,” “AI-generated text detection,” “multi-task learning,” and “joint detection.”

but overoptimization can increase a proxy score while degrading a gold-standard measure (Gao et al., 2023). A detector can achieve useful classification accuracy without its continuous output being calibrated as a measure of style strength or generation quality. Multimodal sarcasm generation may face this risk: ViSP (Wang et al., 2026) uses the DIP sarcasm detector (Wen et al., 2023) as a PPO reward (Schulman et al., 2017). As an exploratory downstream analysis, Section 6 compares detector scores for human authorship with separately obtained, source-blinded LLM ratings of perceived human-likeness.

3 Dataset Construction

Testing for interference requires separate task supervision and evaluation labels for authorship and sarcasm. Here “source” denotes human-versus-AI authorship; “corpus” denotes an input dataset. We combine iSarcasm (Oprea and Magdy, 2020; Abu Farha et al., 2022), CSC (Jang and Frassinelli, 2024; Jang et al., 2023), ToSarcasm (Liang et al., 2022), News Headlines v2 (Misra and Arora, 2023; Misra and Grover, 2021), and the Open Chinese Internet Sarcasm Corpus (Zhu, 2022). None provides both labels. Mapping their annotations to a common binary sarcasm label yields 44,862 pooled instances (approximately 87% English), from which we construct three training sets and DualTest (Appendix A).

Sarcasm classification training set, SarcTrain (10,000, all human). Each language $\times$ label bucket contains 2,500 instances. Three buckets are downsampled; the Chinese sarcastic bucket starts from 2,396 records representing 2,394 unique texts and adds 104 oversampled records. The resulting set is organized into five folds.

MGTD training set, AuthTrain (19,995). From 9,999 language-balanced human seeds, DeepSeek-R1-0528-Qwen3-8B (Ollama tag `deepseek-r1:8b`) (DeepSeek-AI, 2025; Ollama, 2025) produces 9,996 retained outputs (99.97% pairing coverage); Appendix A details quality-control exclusions.

Four language $\times$ style templates follow the balanced seed labels. Authorship and requested-style labels are constructed as separate, approximately balanced factors; no independence is assumed for realized textual features or perceived sarcasm (Appendix A).

Generation training set, GenTrain (2,428; used in §6). From 5,257 eligible positive-sarcasm records in iSarcasm and ToSarcasm, seed-42 sampling selects 2,103 records (1,214 English; 889 Chinese). Chinese oversampling adds 325 records, yielding 1,214 per language. The resulting set contains 374 rewriting, 839 response, and 1,215 comment instances and serves as the instruction-tuning data for SFT (Appendix A).

Diagnostic test set, DualTest (4,664). The human subset combines official iSarcasm (1,400 English) and ToSarcasm (972 Chinese) test sets. Generation quality control retains AI counterparts for 2,292 of 2,372 texts (96.6% coverage). Human texts retain their annotations; for AI texts, the sarcasm label records the prompted condition rather than an independently annotated judgment of perceived sarcasm.

The four source $\times$ sarcasm groups contain 823 human-sarcastic, 1,549 human-non-sarcastic, 779 AI-sarcastic, and 1,513 AI-non-sarcastic instances. Appendix A gives the language $\times$ source $\times$ sarcasm distribution.

4 Dual-Head Detection Model

4.1 Tasks and Architecture

Given a Chinese or English input text x , the model produces two separate, non-mutually-exclusive binary predictions: whether x is labeled sarcastic and whether its authorship source is labeled human or AI. Any combination of the two labels is valid. The tasks share a single DeBERTa-v3-large encoder (He et al., 2023), allowing us to test how incorporating MGTD into joint training changes the representation read by the sarcasm head. This encoder was pretrained on English corpora; §5.1 directly compares it with multilingual-pretrained mBERT / XLM-R on the same bilingual test set. Encoder outputs are aggregated by an intermediate module and then passed to two task-specific classification heads; Figure 1 shows the RCNN configuration.

The RCNN configuration uses BiLSTM + CNN as its intermediate module. A single-layer bidirectional LSTM with hidden size 256 per direction produces a 512-dimensional output at each token position. We concatenate it with the 1,024-dimensional encoder output, then apply three groups of one-dimensional convolutions with kernel widths 3/4/5 and 128 channels each. Global max pooling produces a 384-dimensional vector z .




Figure 1: RCNN dual-head configuration: a shared DeBERTa-v3-large encoder, an RCNN intermediate module (dashed), and two task-specific heads. The comparisons in §5 remove this component in the plain configuration or replace it with task-branch interaction projection, while retaining the same encoder backbone and task definitions (Appendix B).

After dropout, z is passed to the two heads. This module adapts the mBERT-RCNN intermediate architecture of Yang and Ikeda (2024) to a DeBERTa-based dual-head setting.

To test the intermediate layer’s role, §5 compares two variants. The plain model sends [CLS] directly to both heads. The interaction-projection model forms two learned 256-dimensional projections, h_s and h_a , and updates them as $h'_s = h_s + P_{a \rightarrow s}(h_a)$ and $h'_a = h_a + P_{s \rightarrow a}(h_s)$. Each branch contains one vector, so no query-key interaction remains; P is the effective learned value/output projection. The configurations retain the same encoder and task definitions while testing distinct intermediate-representation designs (Appendix B).

4.2 Training Protocol

Training uses a mixed loader that supplies task-homogeneous batches according to a 2:1 MGTD:sarcasm schedule, matching the approximate data-size ratio between AuthTrain and Sarc-

Train. Let θ_{sh} denote parameters updated by both tasks, let $z_s(x)$ and $z_a(x)$ be the representations read by the sarcasm and MGTD heads, and let W_s and W_a denote the two classification heads. The two representations coincide in the plain and RCNN configurations; in interaction projection they are the task-specific h'_s and h'_a defined above. At step t , the loss is computed only for the task corresponding to the current batch B_t :

$$\mathcal{L}(B_t) = \begin{cases} \frac{1}{|B_t|} \sum_{(x,y) \in B_t} \text{CE}_w(W_s z_s(x), y), & B_t \subset \text{SarcTrain}, \\ \frac{1}{|B_t|} \sum_{(x,y) \in B_t} \text{CE}(W_a z_a(x), y), & B_t \subset \text{AuthTrain}, \end{cases} \quad (1)$$

where CE_w is a pre-specified cost-sensitive cross-entropy for sarcasm (weights are provided in Appendix B); because SarcTrain is class-balanced, the weighting emphasizes errors on sarcastic examples rather than correcting training-set imbalance. Softmax is applied to logits only when probabilities are required. The gradient flow follows from the equation: on sarcasm batches, $\partial \mathcal{L} / \partial W_a = 0$ (symmetrically, $\partial \mathcal{L} / \partial W_s = 0$ on MGTD batches), whereas θ_{sh} receives gradients from both types of batches. Unless otherwise stated, binary predictions use a fixed probability threshold of 0.5 rather than a threshold tuned on DualTest.

The encoder uses a lower learning rate than newly added modules. The sarcasm-head rate is fixed across configurations; MGTD-head rates are reported in Appendix B.

Model selection accounts for both tasks. At epoch e , the early-stopping signal is the equally weighted sum of the two validation losses, each normalized by its first-epoch value:

$$\mathcal{L}_{\text{stop}}^{(e)} = \frac{1}{2} \frac{\mathcal{L}_s^{(e)}}{\mathcal{L}_s^{(1)}} + \frac{1}{2} \frac{\mathcal{L}_a^{(e)}}{\mathcal{L}_a^{(1)}}, \quad (2)$$

where $\mathcal{L}_s^{(e)}$ and $\mathcal{L}_a^{(e)}$ are the epoch- e validation losses for the sarcasm and MGTD tasks on their respective validation sets. Normalization reduces the influence of differences in task-specific loss scales and learning dynamics.

Five runs are defined by the SarcTrain validation folds: four folds train and one validates sarcasm, while AuthTrain uses a fixed 90/10 split. Each run reloads the pretrained encoder and randomly initializes new modules and heads. Configurations share fold assignments and nominal seeds, enabling split-matched comparisons but not identical initialization.

DualTest enters no training or validation split. Five-run standard deviations combine fold and initialization variation. Wilson intervals are conditional on a fixed final model and cover test-instance, not training or generator, uncertainty (Table 3; Appendix C.7).

Each subgroup configuration has a final model trained on all applicable data with a fixed epoch count determined from the five validation runs, an independent seed, and no early stopping. Their weights and training duration are not selected using DualTest, and they are not used in five-run aggregates; §6 uses the final RCNN dual-head model (Appendix B).

5 Experiments and Diagnosis

The experiments move from aggregate comparison to dual–single ablation and subgroup diagnosis. We first compare pretrained encoders and intermediate configurations, then measure the change between the evaluated joint- and single-task procedures, and finally partition predictions into source $\times$ sarcasm quadrants. DualTest is used for these descriptive and diagnostic comparisons, but not for training, validation, threshold selection, or epoch selection. Aggregate results are five-run means $\pm$ standard deviations; fixed final models support paired instance-level subgroup comparisons and Wilson intervals.

5.1 Backbone Comparison

We compare dual-head mBERT, XLM-RoBERTa-base, DeBERTa-v3-base, and DeBERTa-v3-large under one evaluation protocol, adding the two intermediate variants on the large encoder. Their pretraining spans multilingual Wikipedia (Devlin et al., 2019), multilingual CommonCrawl (Conneau et al., 2020), and primarily English DeBERTa-v3 pretraining with an ELECTRA-style objective (He et al., 2023; Clark et al., 2020). Positive-class F1 uses the sarcastic class for sarcasm and the human class for MGTD on full DualTest.

Table 1 shows higher mean F1 for the evaluated DeBERTa configurations than for mBERT and XLM-R, with RCNN obtaining the highest means. The intermediate variants also show 34–47% lower observed standard deviations across the five runs. Because backbone scale, pretraining, and intermediate representation vary together, this comparison characterizes the evaluated configurations rather

Model	Sarc. F1	MGTD F1
mBERT dual-head	0.4945 $\pm$.036	0.8743 $\pm$.006
XLM-RoBERTa-base	0.4603 $\pm$.057	0.8321 $\pm$.025
dual-head		
DeBERTa-v3-base	0.5646 $\pm$.056	0.8952 $\pm$.004
head		
DeBERTa-v3-large	0.6126 $\pm$.031	0.9221 $\pm$.010
dual-head		
+ interaction projection	0.6076 $\pm$.020	0.9126 $\pm$.007
+ RCNN	0.6309 $\pm$.020	0.9301 $\pm$.005

Table 1: Descriptive dual-head results across backbones and architectural variants (DualTest; five-run mean $\pm$ standard deviation; bold = highest mean).

Arch.	Task	Dual	Single	Δ
Plain	Sarc.	0.6126 $\pm$.031	0.6734 $\pm$.010	−6.08
Plain	MGTD	0.9221 $\pm$.010	0.9054 $\pm$.027	+1.67
RCNN	Sarc.	0.6309 $\pm$.020	0.6751 $\pm$.022	−4.43
RCNN	MGTD	0.9301 $\pm$.005	0.9253 $\pm$.006	+0.49

Table 2: Dual-head vs. single-head models (mean $\pm$ standard deviation over five fold-based runs; Δ = dual – single, in percentage points).

than isolating one factor. RCNN reaches 0.6309 sarcasm F1 and 0.9301 MGTD F1, establishing a strong shared-encoder diagnostic configuration while leaving a substantial task-level performance gap to explain. Its sarcasm F1 is 0.819 in Chinese and 0.429 in English; prevalence, genre, annotation, and language jointly shape this contrast. Appendix C provides macro-averaged, human-subset, and language-stratified results.

5.2 Dual-Head versus Single-Head Ablation

Within each architecture, dual- and single-head models share the encoder, intermediate structure, sarcasm data, fold assignments, and sarcasm-specific hyperparameters; dual runs additionally receive AuthTrain batches and MGTD updates. Single-head controls traverse the full sarcasm fold each epoch, whereas dual-head models draw sarcasm batches under the 2:1 schedule (§4.2).

Joint training lowers mean sarcasm F1 by 6.08 pp for plain and 4.43 pp for RCNN, whereas the corresponding MGTD changes are +1.67 and +0.49 pp. All ten paired sarcasm differences are negative (−0.8 to −8.6 pp), so the transfer cost recurs across every split-matched run and both architectures. This repeated asymmetry identifies sarcasm as the task bearing the aggregate cost and motivates locating that cost by source and class. The smaller dual–single gap under RCNN further motivates the intermediate-architecture di-

agnosis below; capacity and compute differences are recorded in Limitations.

5.3 Four-Quadrant Diagnosis: Class-Conditional Negative Transfer across Intermediate Architectures

Aggregate F1 hides whether errors depend jointly on source and sarcasm class. We therefore partition DualTest by source $\times$ sarcasm. Human labels are dataset annotations and AI labels record prompted conditions; under these operational task labels, single-class quadrants use recall for sarcastic subsets and false-positive rate for non-sarcastic subsets.

5.3.1 Three Architectures and Single-Head Controls

The three dual-head variants show how the subgroup pattern changes with intermediate representation. Plain and RCNN single-head models provide same-architecture controls for the joint-training comparison. Because these controls retain the same all-human sarcasm data, the comparison shows that the shift emerges under the evaluated joint-training procedure rather than from adding AI texts to SarcTrain. Table 3 reports sarcasm predictions; corresponding MGTD results are in Appendix C.

The plain dual-head model predicts sarcasm for 86.3% of AI $\times$ non-sarcastic instances. The observed rates are 40.9% for interaction projection and 34.1% for RCNN, while both same-architecture single-head controls yield 19.3%. Thus, the observed severity is more than 45 pp lower in the two evaluated intermediate configurations than in the plain dual-head configuration.

On the same texts, plain dual versus single gives 1,021/7 discordant pairs ($p < .001$). The same-architecture single-head model receives the same all-human sarcasm supervision but not the joint MGTD procedure, tying the large increase to the evaluated joint-training procedure. RCNN’s same-architecture comparison shows the same direction with a smaller gap; interaction projection has no single-head counterpart and contributes complementary configuration-level evidence.

No analogous extreme subgroup error appears in the MGTD head: its false-positive rate is 6.1% for plain and 6.6% for RCNN. Within AI texts, sarcasm AUC remains 0.82–0.88 while joint training raises score levels and largely preserves ordering. The failure is therefore not a collapse of within-AI

ranking, but a source-conditioned upward shift that produces class-conditional negative transfer on the non-sarcastic prompted subset (Appendix C.7).

The matched plain comparison makes the concentration especially clear: its non-sarcastic false-positive rate changes by 4.2 pp on human text but by 67.0 pp on AI-generated text. The aggregate F1 decrease is therefore driven by a sharply localized source $\times$ class effect rather than a uniform loss across DualTest.

5.3.2 Mechanism Hypothesis and Representation-Level Diagnostics

We hypothesize that MGTD-shaped representations expose source-correlated cues to a sarcasm head that receives no sarcasm supervision on AI text. Source is highly decodable in both dual-head (AUC 0.986) and same-architecture single-head (0.974) representations, so its presence alone is not the explanation. In the dual-head model, however, the MGTD-head AI logit correlates with the sarcasm-head logit on the target subgroup ($\rho = 0.519$), while both MGTD-head and probe directions align with the sarcasm direction in whitened geometry.

Representation interventions further characterize the effective variation. LEACE (Belrose et al., 2023) reduces measured linear source decodability from AUC 0.982 to 0.589 and shifts the non-sarcastic false-positive rates from 84.8%/32.8% on AI/human text to 68.7%/49.5%. Removing a high-variance direction orthogonal to the estimated linear source direction yields 66.7%/54.1% while preserving measured source decodability; ten sampled random controls show no comparable source-specific shift. Under LEACE, the unweighted mean across the two source groups remains nearly unchanged (58.8% to 59.1%), showing error redistribution rather than overall mitigation. Together, these results support a bounded source-correlated shortcut hypothesis and show prediction sensitivity to high-variance directions without isolating a unique causal feature (Appendix E).

This evidence chain goes beyond probe accuracy alone: task-head logit correlation connects source-related behavior to the decision, targeted interventions change how errors are allocated across sources, and the orthogonal and random controls constrain the interpretation. The result is a representation-level diagnosis of score sensitivity rather than a causal localization inferred from output labels alone.

Model	H sarc. rec.	H non-sarc. FP	AI sarc. rec.	AI non-sarc. FP
Plain dual-head	67.1%	32.9%	98.2%	86.3% [84.5, 88.0]
Interaction-projection dual-head	71.9%	37.6%	94.4%	40.9% [38.5, 43.4]
RCNN dual-head	80.8%	40.5%	89.3%	34.1% [31.8, 36.5]
Plain single-head (sarcasm only)	67.8%	28.7%	76.6%	19.3% [17.4, 21.4]
RCNN single-head (sarcasm only)	66.2%	28.5%	79.1%	19.3% [17.4, 21.4]

Table 3: Five sarcasm classifiers on four source $\times$ sarcasm quadrants (recall / false-positive rate, FP). Human labels are annotations; AI labels record prompted conditions. Brackets give model-conditional Wilson 95% intervals; full intervals and paired tests are in Appendix C.7.

5.3.3 Cross-Generator Replication

We regenerate only DualTest’s AI side with Qwen3.5-9B (Ollama tag qwen3.5:9b) (Qwen Team, 2026), retaining the human texts and trained detectors. Among 1,550 non-sarcastic-condition outputs, plain dual rises from 86.3% to 91.3% and its same-architecture single falls from 19.3% to 7.9%, increasing the gap from 67.0 to 83.4 pp. RCNN and interaction projection remain lower at 22.1% and 29.5%. Thus, both generators preserve the large plain dual–single separation and the lower rates of the two intermediate configurations. The second generator reproduces the gap direction with a different magnitude (Appendix C.7).

6 Exploratory Comparison of Detector Scores and LLM-Judge Ratings

As a downstream implication, we test whether source-classification probability tracks perceived human-likeness and whether sarcasm probability tracks LLM ratings. Human-authorship probability and perceived human-likeness are distinct constructs; their comparison directly tests whether a classifier score can function as a generation-evaluation proxy. We compare 600 outputs from Qwen3-8B (Yang et al., 2025) before and after QLoRA fine-tuning on GenTrain. The RCNN detector and three LLM judges (DeepSeek-AI, 2024; Comanici et al., 2025; Singh et al., 2025) score outputs from 180 Chinese topic-comment and 120 English rewriting prompts; the judges are not shown the ground-truth source, and language is coupled with generation task.

Agreement varies by dimension and language. In Chinese, SFT raises detector sarcasm from 0.780 to 0.895 but lowers mean judge rating from 4.01 to 3.71; sarcasm correlation is -0.020 in Chinese and $+0.338$ in English. SFT outputs also receive higher human-authorship probabilities but lower perceived human-likeness ratings, yielding a pooled correlation of -0.212 . Within-condition

correlations show that this negative association is driven mainly by the reversal between base and SFT conditions.

This condition-level reversal is directly relevant to generation evaluation: fine-tuning moves the detector scores in the apparently favorable direction while moving the external LLM ratings downward. A detector can therefore provide a useful classification boundary without its continuous score preserving the ordering required of a style or quality measure.

Inter-judge agreement is low to moderate at the instance level, but all six language $\times$ judge mean comparisons rate SFT below the base outputs, five significantly after FDR correction. A plain-detector control changes the reported correlations by at most 0.02. The consistent condition-level directions therefore show that detector probabilities require external validation as generation-quality proxies or rewards. Full protocols, tests, and representative cases are reported in Appendices C–D.

7 Conclusion

Joint training can systematically distort the sarcasm decision even when source and sarcasm labels do not conflict. Relative to a same-architecture single-head control, the plain dual-head model raises the AI non-sarcastic false-positive rate by 67.0 pp; the effect holds in both languages and replicates with a second generator, while RCNN and interaction projection show substantially lower rates. Representation analyses connect the shift to source-correlated, high-variance structure and controlled error redistribution. Detector–judge reversals further show why classification probabilities require external validation as style or quality measures. Aggregate metrics should therefore be accompanied by source $\times$ class analysis.

Limitations

DualTest combines human sarcasm annotations with instruction-conditioned labels for AI texts, so AI-side false positives are defined relative to requested style rather than independent perceived-sarcasm annotation. Main-generator failures leave 80 unpaired human texts. Seven train–test matches correspond to three unique AI texts; removing them changes all core target rates by less than 0.1 pp. Some SarcTrain duplicates span validation folds and may affect checkpoint selection, but DualTest remains independent.

The core dual–single comparisons identify the complete joint-training procedure: architecture, sarcasm data, folds, sarcasm-head learning rate, and class weights are shared, while MGTD batches, task scheduling, shared updates, and joint early stopping enter together. Compute-matched controls could further decompose this effect. Cross-architecture results are configuration-level because capacity and optimization differ, and the cost-sensitive loss can affect absolute rates. MGTD training uses one generator; the second-generator test establishes directional replication beyond it.

The fixed-model linear analyses connect source-related structure to behavior but do not isolate a unique causal feature or training-seed distribution. The detector–judge study uses LLM rather than human ratings and couples language with generation task. These boundaries motivate independent style annotation, matched-budget controls, cross-seed mechanism tests, and broader generator and human-evaluation studies.

Ethical considerations

All human-authored texts come from publicly released research datasets: iSarcasm, CSC, ToSarcasm, News Headlines, and the Open Chinese Internet Sarcasm Corpus. Explicit repository licenses were available for iSarcasm and CSC; for sources without an identified redistribution license, we do not redistribute the human-authored text (Appendix A). We collect no new user data. We did not conduct an additional personally identifiable information or offensive-content audit beyond the preparation reflected in the released corpora.

All AI texts were generated locally using publicly released open-weight models (DeepSeek-R1-0528-Qwen3-8B and Qwen3.5-9B). They are used only for training and evaluation and are labeled

as AI-generated. The study recruits no new human participants, but analyzes human-authored texts from previously released datasets. The automated evaluation transmits each evaluation input, generated output, and rating instructions to the DeepSeek-V3, Gemini-2.5-Flash-Lite, and GPT-5-nano APIs; it does not transmit dataset metadata or ground-truth source labels.

The detectors are developed for research diagnosis. Their outputs should not be used as evidence of authorship or misconduct at the individual-text level without human review and domain-specific validation, nor deployed for content moderation or automated sanctions without evaluating out-of-distribution behavior. The class-conditional failure diagnosed here is an example of systematic subgroup error affecting a particular type of generated content.

AI-assisted tools were used for grammar correction, wording suggestions, and formatting. They did not generate the reported experimental outputs or statistics; the authors determined the study design, analysis, and conclusions and take full responsibility for the work.

References

- Ibrahim Abu Farha, Silviu Vlad Oprea, Steven Wilson, and Walid Magdy. 2022. [SemEval-2022 task 6: iSarcasmEval, intended sarcasm detection in English and Arabic](#). In *Proceedings of the 16th International Workshop on Semantic Evaluation (SemEval-2022)*.
- Nora Belrose, David Schneider-Joseph, Shauli Ravfogel, Ryan Cotterell, Edward Raff, and Stella Biderman. 2023. [LEACE: Perfect linear concept erasure in closed form](#). In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Yoav Benjamini and Yosef Hochberg. 1995. [Controlling the false discovery rate: A practical and powerful approach to multiple testing](#). *Journal of the Royal Statistical Society: Series B (Methodological)*, 57(1):289–300.
- Jane Bevendorff, Matti Wiegmann, Emmelie Richter, Martin Potthast, and Benno Stein. 2025. [The two paradigms of LLM detection: Authorship attribution vs. authorship verification](#). In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 3762–3787.
- Silvia Casola, Simona Frenda, Soda Marem Lo, Erhan Sezerer, Antonio Uva, Valerio Basile, Cristina Bosco, Alessandro Pedrani, Chiara Rubagotti, Viviana Patti, and Davide Bernardi. 2024. [MultiPICO: Multilingual perspectivist irony corpus](#). In *Proceedings of the 62nd Annual Meeting of the Association*

807	for Computational Linguistics (ACL), pages 16008–	862
808	16021.	863
809	Junyi Chai, Shenyu Lu, and Xiaoqian Wang. 2025.	864
810	Identifying and mitigating spurious correlation	865
811	in multi-task learning. In <i>Proceedings of the</i>	
812	<i>IEEE/CVF Conference on Computer Vision and Pat-</i>	
813	<i>tern Recognition (CVPR)</i> , pages 25698–25707.	
814	Kevin Clark, Minh-Thang Luong, Quoc V. Le, and	
815	Christopher D. Manning. 2020. ELECTRA: Pre-	
816	training text encoders as discriminators rather than	
817	generators. In <i>Proceedings of the 8th International</i>	
818	<i>Conference on Learning Representations (ICLR)</i> .	
819	Gheorghe Comanici et al. 2025. Gemini 2.5: Pushing	
820	the frontier with advanced reasoning, multimodality,	
821	long context, and next generation agentic capabili-	
822	ties. <i>arXiv preprint arXiv:2507.06261</i> .	
823	Rhys Compton, Lily Zhang, Aahlad Puli, and Rajesh	
824	Ranganath. 2023. When more is less: Incorporating	
825	additional datasets can hurt performance by introduc-	
826	ing spurious correlations. In <i>Proceedings of the 8th</i>	
827	<i>Machine Learning for Healthcare Conference</i> , vol-	
828	ume 219 of <i>Proceedings of Machine Learning Re-</i>	
829	<i>search</i> , pages 110–127. PMLR.	
830	Alexis Conneau, Kartikay Khandelwal, Naman Goyal,	
831	Vishrav Chaudhary, Guillaume Wenzek, Francisco	
832	Guzmán, Edouard Grave, Myle Ott, Luke Zettle-	
833	moyer, and Veselin Stoyanov. 2020. Unsupervised	
834	cross-lingual representation learning at scale. In <i>Pro-</i>	
835	<i>ceedings of the 58th Annual Meeting of the Associa-</i>	
836	<i>tion for Computational Linguistics (ACL)</i> .	
837	DeepSeek-AI. 2024. DeepSeek-V3 technical report.	
838	<i>arXiv preprint arXiv:2412.19437</i> .	
839	DeepSeek-AI. 2025. DeepSeek-R1: Incentivizing rea-	
840	soning capability in LLMs via reinforcement learn-	
841	ing. <i>arXiv preprint arXiv:2501.12948</i> .	
842	Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and	
843	Luke Zettlemoyer. 2023. QLoRA: Efficient finetun-	
844	ing of quantized LLMs. In <i>Advances in Neural In-</i>	
845	<i>formation Processing Systems (NeurIPS)</i> .	
846	Jacob Devlin, Ming-Wei Chang, Kenton Lee, and	
847	Kristina Toutanova. 2019. BERT: Pre-training of	
848	deep bidirectional transformers for language under-	
849	standing. In <i>Proceedings of the 2019 Conference of</i>	
850	<i>the North American Chapter of the Association for</i>	
851	<i>Computational Linguistics: Human Language Tech-</i>	
852	<i>nologies (NAACL-HLT)</i> .	
853	Xiyang Du, Dou Hu, Jin Zhi, Lianxin Jiang, and Xi-	
854	aofeng Shi. 2022. PALI-NLP at SemEval-2022 task	
855	6: iSarcasmEval- fine-tuning the pre-trained model	
856	for detecting intended sarcasm. In <i>Proceedings of</i>	
857	<i>the 16th International Workshop on Semantic Evalua-</i>	
858	<i>tion (SemEval-2022)</i> , pages 815–819.	
859	Liam Dugan, Alyssa Hwang, Filip Trhlik, Andrew	
860	Zhu, Josh Magnus Ludan, Hainiu Xu, Daphne Ip-	
861	polito, and Chris Callison-Burch. 2024. RAID: A	
	shared benchmark for robust evaluation of machine-	862
	generated text detectors. In <i>Proceedings of the</i>	863
	<i>62nd Annual Meeting of the Association for Compu-</i>	864
	<i>tational Linguistics (ACL)</i> , pages 12463–12492.	865
	Yanai Elazar, Shauli Ravfogel, Alon Jacovi, and Yoav	866
	Goldberg. 2021. Amnesic probing: Behavioral ex-	867
	planation with amnesic counterfactuals. <i>Transac-</i>	868
	<i>tions of the Association for Computational Linguis-</i>	869
	<i>tics (TACL)</i> , 9:160–175.	870
	Leo Gao, John Schulman, and Jacob Hilton. 2023. Scal-	871
	ing laws for reward model overoptimization. In <i>Pro-</i>	872
	<i>ceedings of the 40th International Conference on</i>	873
	<i>Machine Learning (ICML)</i> , pages 10835–10866.	874
	Robert Geirhos, Jörn-Henrik Jacobsen, Claudio	875
	Michaelis, Richard Zemel, Wieland Brendel,	876
	Matthias Bethge, and Felix A. Wichmann. 2020.	877
	Shortcut learning in deep neural networks. <i>Nature</i>	878
	<i>Machine Intelligence</i> , 2:665–673.	879
	Aniruddha Ghosh and Tony Veale. 2016. Fracking	880
	sarcasm using neural network. In <i>Proceedings of</i>	881
	<i>the 7th Workshop on Computational Approaches to</i>	882
	<i>Subjectivity, Sentiment and Social Media Analysis</i>	883
	<i>(WASSA)</i> .	884
	Biyang Guo, Xin Zhang, Ziyuan Wang, Minqi Jiang,	885
	Jinran Nie, Yuxuan Ding, Jianwei Yue, and Yupeng	886
	Wu. 2023. How close is ChatGPT to human experts?	887
	comparison corpus, evaluation, and detection. <i>arXiv</i>	888
	<i>preprint arXiv:2301.07597</i> .	889
	Han He and Jinho D. Choi. 2021. The stem cell hy-	890
	pothesis: Dilemma behind multi-task learning with	891
	transformer encoders. In <i>Proceedings of the 2021</i>	892
	<i>Conference on Empirical Methods in Natural Lan-</i>	893
	<i>guage Processing (EMNLP)</i> .	894
	Pengcheng He, Jianfeng Gao, and Weizhu Chen. 2023.	895
	DeBERTaV3: Improving DeBERTa using ELEC-	896
	TRA-style pre-training with gradient-disentangled	897
	embedding sharing. In <i>International Conference on</i>	898
	<i>Learning Representations (ICLR)</i> .	899
	John Hewitt and Percy Liang. 2019. Designing and in-	900
	terpreting probes with control tasks. In <i>Proceedings</i>	901
	<i>of the 2019 Conference on Empirical Methods in Nat-</i>	902
	<i>ural Language Processing and the 9th International</i>	903
	<i>Joint Conference on Natural Language Processing</i>	904
	<i>(EMNLP-IJCNLP)</i> .	905
	Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and	906
	Yejin Choi. 2020. The curious case of neural text de-	907
	generation. In <i>International Conference on Learn-</i>	908
	<i>ing Representations (ICLR)</i> .	909
	Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan	910
	Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and	911
	Weizhu Chen. 2022a. LoRA: Low-rank adaptation	912
	of large language models. In <i>International Confer-</i>	913
	<i>ence on Learning Representations (ICLR)</i> .	914

915	Ziniu Hu, Zhe Zhao, Xinyang Yi, Tiansheng Yao, Lichan Hong, Yizhou Sun, and Ed H. Chi. 2022b. Improving multi-task generalization via regularizing spurious correlation . In <i>Advances in Neural Information Processing Systems (NeurIPS)</i> .	971	Rishabh Misra and Jigyasa Grover. 2021. <i>Sculpting Data for ML: The first act of Machine Learning</i> . Independently published.	972
916		973		974
917			Eric Mitchell, Yoonho Lee, Alexander Khazatsky, Christopher D. Manning, and Chelsea Finn. 2023. DetectGPT: Zero-shot machine-generated text detection using probability curvature . In <i>Proceedings of the 40th International Conference on Machine Learning (ICML)</i> .	975
918				976
919				977
920	Hyewon Jang, Bettina Braun, and Diego Frassinelli. 2023. Intended and perceived sarcasm between close friends: What triggers sarcasm and what gets conveyed? In <i>Proceedings of the Annual Meeting of the Cognitive Science Society</i> , volume 45.	921		978
922		922		979
923		923		
924		924		
925	Hyewon Jang and Diego Frassinelli. 2024. Generalizable sarcasm detection is just around the corner, of course! In <i>Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT)</i> , pages 4238–4249.	925		
926		926		
927		927		
928		928		
929		929		
930		930		
931	Rabeeh Karimi Mahabadi, Sebastian Ruder, Mostafa Dehghani, and James Henderson. 2021. Parameter-efficient multi-task fine-tuning for transformers via shared hypernetworks . In <i>Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics (ACL)</i> .	931		
932		932		
933		933		
934		934		
935		935		
936		936		
937	Akshay Khatrri and Pranav P. 2020. Sarcasm detection in tweets with BERT and GloVe embeddings . In <i>Proceedings of the Second Workshop on Figurative Language Processing (FigLang)</i> .	937		
938		938		
939		939		
940		940		
941	Taku Kudo and John Richardson. 2018. SentencePiece: A simple and language independent subword tokenizer and detokenizer for neural text processing . In <i>Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing: System Demonstrations</i> , pages 66–71.	941		
942		942		
943		943		
944		944		
945		945		
946		946		
947	Siwei Lai, Liheng Xu, Kang Liu, and Jun Zhao. 2015. Recurrent convolutional neural networks for text classification . In <i>Proceedings of the AAAI Conference on Artificial Intelligence</i> , volume 29, pages 2267–2273.	947		
948		948		
949		949		
950		950		
951		951		
952	Chong Li, Shaonan Wang, Yunhao Zhang, Jiajun Zhang, and Chengqing Zong. 2023. Interpreting and exploiting functional specialization in multi-head attention under multi-task learning . In <i>Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)</i> , pages 16460–16476.	952		
953		953		
954		954		
955		955		
956		956		
957		957		
958		958		
959	Bin Liang, Zijie Lin, Bing Qin, and Ruifeng Xu. 2022. Topic-oriented sarcasm detection: New task, new dataset and new method . In <i>Proceedings of the 21st Chinese National Conference on Computational Linguistics (CCL)</i> , pages 557–568.	959		
960		960		
961		961		
962		962		
963		963		
964	Bo Liu, Xingchao Liu, Xiaojie Jin, Peter Stone, and Qiang Liu. 2021. Conflict-averse gradient descent for multi-task learning . In <i>Advances in Neural Information Processing Systems (NeurIPS)</i> .	964		
965		965		
966		966		
967		967		
968	Rishabh Misra and Prahal Arora. 2023. Sarcasm detection using news headlines dataset . <i>AI Open</i> , 4:13–18.	968		
969		969		
970		970		
				971
				972
				973
				974
				975
				976
				977
				978
				979
				980
				981
				982
				983
				984
				985
				986
				987
				988
				989
				990
				991
				992
				993
				994
				995
				996
				997
				998
				999
				1000
				1001
				1002
				1003
				1004
				1005
				1006
				1007
				1008
				1009
				1010
				1011
				1012
				1013
				1014
				1015
				1016
				1017
				1018
				1019
				1020
				1021
				1022
				1023
				1024
				1025

1026	Aaditya Singh et al. 2025. OpenAI GPT-5 system card . <i>arXiv preprint arXiv:2601.03267</i> .	1081
1027		1082
1028	Asa Cooper Stickland and Iain Murray. 2019. BERT and PALs: Projected attention layers for efficient adaptation in multi-task learning . In <i>Proceedings of the 36th International Conference on Machine Learning (ICML)</i> .	1083
1029		1084
1030		1085
1031		1086
1032		1087
1033	Yi Tay, Luu Anh Tuan, Siu Cheung Hui, and Jian Su. 2018. Reasoning with sarcasm by reading in-between . In <i>Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (ACL)</i> .	1088
1034		1089
1035		1090
1036		1091
1037		1092
1038	Bertie Vidgen, Alex Harris, Dong Nguyen, Rebekah Tromble, Scott Hale, and Helen Margetts. 2019. Challenges and frontiers in abusive content detection . In <i>Proceedings of the Third Workshop on Abusive Language Online</i> , pages 80–93.	1093
1039		1094
1040		1095
1041		1096
1042		
1043	Changli Wang, Fang Yin, Jiafeng Liu, and Rui Wu. 2026. ViSP: A PPO-enhanced framework for multi-modal sarcasm generation with contrastive learning . <i>Neurocomputing</i> , 678:133185.	1097
1044		1098
1045		1099
1046		1100
1047	Yuxia Wang, Jonibek Mansurov, Petar Ivanov, Jinyan Su, Artem Shelmanov, Akim Tsvigun, Chenxi Whitehouse, Osama Mohammed Afzal, Tarek Mahmoud, Toru Sasaki, Thomas Arnold, Alham Fikri Aji, Nizar Habash, Iryna Gurevych, and Preslav Nakov. 2024. M4: Multi-generator, multi-domain, and multi-lingual black-box machine-generated text detection . In <i>Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 1369–1407.	1101
1048		
1049		1102
1050		1103
1051		1104
1052		1105
1053		1106
1054		1107
1055		
1056		1108
1057		1109
1058	Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc V. Le, and Denny Zhou. 2022. Chain-of-thought prompting elicits reasoning in large language models . In <i>Advances in Neural Information Processing Systems (NeurIPS)</i> , pages 24824–24837.	1110
1059		1111
1060		1112
1061		1113
1062		1114
1063		
1064	Changsong Wen, Guoli Jia, and Jufeng Yang. 2023. DIP: Dual incongruity perceiving network for sarcasm detection . In <i>Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)</i> , pages 2540–2550.	1115
1065		1116
1066		1117
1067		1118
1068		1119
1069	Lianwei Wu, Yuan Rao, Haolin Jin, Ambreen Nazir, and Ling Sun. 2019. Different absorption from the same sharing: Sifted multi-task learning for fake news detection . In <i>Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)</i> .	1120
1070		1121
1071		1122
1072		1123
1073		1124
1074		
1075		1125
1076		1126
1077	Sen Wu, Hongyang R. Zhang, and Christopher Ré. 2020. Understanding and improving information transfer in multi-task learning . In <i>International Conference on Learning Representations (ICLR)</i> .	1127
1078		1128
1079		1129
1080		1130
		1131
		1132
		1133
		1134
	An Yang et al. 2025. Qwen3 technical report . <i>arXiv preprint arXiv:2505.09388</i> .	
	Linshuo Yang and Daisuke Ikeda. 2024. The influence of linguistic attribute differences in multilingual datasets on sarcasm detection . <i>International Journal of Service and Knowledge Management</i> , 8(2):1–15.	
	Wenqian Ye, Luyang Jiang, Eric Xie, Guangtao Zheng, Yunsheng Ma, Xu Cao, Dongliang Guo, Daiqing Qi, Zeyu He, Yijun Tian, Christopher W. Porter, Megan Coffee, Zhe Zeng, Sheng Li, Ziran Wang, Ting-Hao (Kenneth) Huang, James Matthew Rehg, Henry Kautz, and Aidong Zhang. 2026. The clever Hans mirage: A comprehensive survey on spurious correlations in machine learning . <i>Transactions on Machine Learning Research</i> .	
	Tianhe Yu, Saurabh Kumar, Abhishek Gupta, Sergey Levine, Karol Hausman, and Chelsea Finn. 2020. Gradient surgery for multi-task learning . In <i>Advances in Neural Information Processing Systems (NeurIPS)</i> .	
	Tan Yue, Xuzhao Shi, Rui Mao, Zonghai Hu, and Erik Cambria. 2024. SarcNet: A multilingual multimodal sarcasm detection dataset . In <i>Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING)</i> , pages 14325–14335.	
	Zhihan Zhang, Wenhao Yu, Mengxia Yu, Zhichun Guo, and Meng Jiang. 2023. A survey of multi-task learning in natural language processing: Regarding task relatedness and training methods . In <i>Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics (EACL)</i> , pages 943–956.	
	Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. Judging LLM-as-a-judge with MT-Bench and Chatbot Arena . In <i>Advances in Neural Information Processing Systems (NeurIPS)</i> .	
	Yizhang Zhu. 2022. Open Chinese internet sarcasm corpus construction: An approach . <i>Frontiers in Computing and Intelligent Systems</i> , 2(1):7–9.	

A Detailed Dataset Statistics and Quality Checks

A.1 Corpus details. The pooled corpus collection contains 44,862 instances in total (86.9% English / 13.1% Chinese). Chinese–English 1:1 balance in SarcTrain is achieved through bucketed sampling. The two English buckets and the Chinese non-sarcastic bucket contain more than 2,500 available instances and are randomly downsampled within each bucket. The Chinese sarcastic

Corpus	Lang.	Inst.	Use
iSarcasm (SemEval-2022 T6)	English	3,467	Sarc. train + SFT
CSC	English	6,878	Sarc. train
ToSarcasm	zh 3,892 / en 6	3,898	Sarc. train + SFT
News Headlines v2	English	28,619	Sarc. train
Open Chinese Internet Sarcasm Corpus	Chinese	2,000	Sarc. train

Table 4: Input corpora. Genre / annotation: iSarcasm — tweets, speaker self-annotation; CSC — conversational responses; ToSarcasm — topic comments, third-party annotation; News Headlines — distant supervision (The Onion / HuffPost); Open Chinese Corpus — Weibo, third-party annotation.

bucket contains 2,396 source records representing 2,394 unique texts; 104 records are added by oversampling to reach 2,500. Including source-level duplicate text, the final bucket’s text-duplication rate is 4.24% (A.3).

Licensing and release scope. The iSarcasm and CSC repositories state MIT licenses. We did not identify explicit redistribution licenses in the releases we accessed for ToSarcasm, News Headlines v2, or the Open Chinese Internet Sarcasm Corpus. We therefore do not redistribute human-authored text from those sources. Planned research artifacts are limited to eligible AI-generated text, labels and metadata, and reconstruction scripts that require users to obtain the human-authored corpora from their original sources.

Lang.	Source	Sarc.	Non-sarc.
English	Human	200	1,200
English	AI	192	1,174
Chinese	Human	623	349
Chinese	AI	587	339

Table 5: Language $\times$ source $\times$ sarcasm distribution of DualTest. Human labels are dataset annotations; AI labels record prompted conditions.

A.2 Three-dimensional distribution of Dual-Test.

A.3 Data quality checks. Automated checks confirm language and label balance across the five SarcTrain folds. AuthTrain contains 9,996 complete human–AI pairs and three unmatched human instances. Exact-match screening finds two

SarcTrain instances (0.02%) and five AuthTrain instances (0.025%) that also occur on the AI side of DualTest; overlap on the human subset is zero. The seven records correspond to three unique DualTest texts. Removing those three texts at evaluation time changes every core AI $\times$ non-sarcastic rate by less than 0.1 pp; this sensitivity check bounds their direct contribution to the reported test rates but does not remove their training-side influence.

SarcTrain’s text-duplication rate is 0.34% for English / 2.16% for Chinese. Chinese duplication is concentrated in the sarcastic bucket (4.24%; non-sarcastic bucket: 0.08%) and results from the oversampling used to fill that bucket (A.1). Because oversampling precedes fold assignment, 86 of the 104 added records have a source counterpart in another fold. Under the raw-text counting convention used for these duplication rates, 116 repeated-text groups involving 241 records are present when oversampled copies are included, and 98 groups span folds. Excluding the added copies and normalizing whitespace and case leaves 19 naturally duplicated groups involving 46 records, of which 16 span folds. These relationships are relevant to fold-based checkpoint selection but do not place the affected training or validation records into DualTest. Post-processing also checks for the explicit chain-of-thought markers and reasoning prefaces described in A.4 (Wei et al., 2022).

A.4 AI generation protocol. All AI texts are produced by locally deployed generation models through the Ollama API (DeepSeek-R1-0528-Qwen3-8B under tag `deepseek-r1:8b` for the main results; Qwen3.5-9B under tag `qwen3.5:9b` for cross-generator replication), using temperature = 0.8, top- p = 0.95 (Holtzman et al., 2020), and a generation limit of 500 tokens. Generation instructions use four templates defined by language $\times$ style, with style determined by the seed instance’s sarcasm label. The templates constrain output style and length (20–80 Chinese characters; 10–50 English words). This constraint is intended to reduce, but does not eliminate, source differences in length; Appendix E separately repeats the representation correlation analysis within length tertiles. The complete Chinese sarcastic template used in the experiment is reproduced verbatim below:

【任务】生成一条讽刺风格的评论
 【主题】{seed}
 【要求】
 - 直接给出简短的评论，不要解释思路或分析
 - 风格要讽刺、幽默、夸张

- 表达与主题相关的观点即可, 不必追求严谨
 - 20-80 字为佳
- 【评论】

Its pinyin transliteration is:

[rènwù] shēngchéng yī tiáo fěngcì
fěnggé de pínglùn
[zhǔtí] {seed}
[yāoqiú]
- zhíjiē gěichū jiǎnduǎn de pínglùn,
bùyào jiěshì sīlù huò fēnxī
- fěnggé yào fěngcì, yōumò, kuāzhāng
- biǎodá yǔ zhǔtí xiāngguān de guāndiǎn
jíkě, bùbì zhuīqiú yánjīn
- 20-80 zì wéi jiā
[pínglùn]

Its English translation is:

[Task] Generate a sarcastic-style
comment
[Topic] {seed}
[Requirements]
- Directly provide a short comment;
do not explain your reasoning or
analysis
- The style should be sarcastic,
humorous, and exaggerated
- It is sufficient to express a
viewpoint related to the topic;
strict rigor is not required
- Preferably 20--80 Chinese
characters
[Comment]

The Chinese non-sarcastic template changes only the style line to “The style should be normal and natural.” The English templates follow the same structure, using either “sarcastic, humorous, exaggerated” or “normal, natural,” with a length limit of 10–50 words.

Post-processing removes `<think>...</think>` segments, reasoning prefaces such as “好的” (*hǎo de* ‘okay’), “嗯” (*ní* ‘hmm’), “Okay,” “Let me,” and “First,” and Markdown markup. Outputs that remain overlength or cyclic after repeated generation attempts are discarded. AuthTrain retains 9,996 AI counterparts from 9,999 seeds; DualTest retains 2,292 counterparts from 2,372 seeds.

The Qwen3.5-9B replication has no generation failures, yielding 1,550 AI × non-sarcastic instances rather than 1,513. Generator-specific rates therefore use each retained full output set rather than an identical seed intersection. For this run, Ollama thinking mode is disabled (`think = false`) to prevent the reasoning process from

exhausting the output budget. The same post-processing removes any inline chain-of-thought from DeepSeek-R1-0528-Qwen3-8B.

A.5 SFT training and generation-evaluation set.

Eligibility filtering retains 5,257 positive-sarcasm records suitable for generation tasks (2,972 English; 2,285 Chinese). A seed-42 shuffle and 40% sampling select 2,103 records (1,214 English; 889 Chinese); language balancing then adds 325 oversampled Chinese records, producing GenTrain (2,428; 1,214 per language). By prompt type, it comprises 374 rewriting, 839 response, and 1,215 comment instances. Generation evaluation uses an independent set of 300 prompts (180 Chinese and 120 English): the Chinese task is to generate a sarcastic comment on a given topic, and the English task is to rewrite a normal sentence as a sarcastic one. The detector scoring and three source-blinded LLM evaluations in §6 are both conducted on this set.

B Hyperparameters and Training Details

Group	Enc.	Inter.	Sarc.	MGTD
Plain / backbone dual	5e-7	—	5e-6	1e-5
RCNN dual	5e-7	5e-6	5e-6	5e-6
Interaction-projection dual	5e-7	5e-6	5e-6	5e-6
Plain single-task	5e-7	—	5e-6	1e-5
RCNN single-task	5e-7	5e-6	5e-6	1e-5

Table 6: Discriminative learning rates for the evaluated dual- and single-task configurations.

B.1 Discriminative learning rates. The sarcasm-head learning rate is fixed at 5e-6 across configurations. The MGTD-head rate is 1e-5 for the plain dual and single-task models, 5e-6 for the two intermediate dual-head variants, and 1e-5 for the RCNN MGTD single-task control. Thus, the core sarcasm dual-single comparisons retain the same sarcasm-specific rates within each architecture, whereas the RCNN MGTD comparison and cross-configuration results can also reflect optimization settings.

B.2 Shared training hyperparameters. Training uses AdamW with default PyTorch betas (0.9, 0.999) and $\epsilon = 10^{-8}$, weight decay 0.01, and parameter-group learning rates in B.1. A LambdaLR schedule linearly warms up for 200 steps and then keeps the learning rate constant. Training uses at most 10 epochs, early-stopping patience

5, and gradient clipping 1.0. The physical batch size is 4 with gradient accumulation 4, giving an effective batch size of 16. The MGTD:sarcasm sampling ratio is 2:1, and sarcasm class weights are [1.0, 2.5]. Both choices were specified before testing. Because SarcTrain is balanced, the class weights implement cost-sensitive emphasis rather than imbalance correction. The normalized joint validation loss is $0.5 \times [\text{sarcasm validation loss} / \text{first-epoch value}] + 0.5 \times [\text{MGTD validation loss} / \text{first-epoch value}]$. The single-head controls traverse the full sarcasm training fold each epoch, whereas the dual-head models draw sarcasm batches under the 2:1 task-mixing schedule.

Detector inputs use a maximum sequence length of 512 tokens. In the RCNN module, a width- k convolution uses padding $\lfloor k/2 \rfloor$ followed by ReLU; global max pooling is applied directly over the full convolution output without an additional padding mask.

All fold-based runs start from seed 42 plus the fold index and share the same fold assignments and nominal seed scheme across architectures and ablations. MGTD data use a fixed 90/10 holdout in every fold. For each final-model configuration, the fixed epoch count is the mode of its five best epoch counts plus one, capped at 10 epochs; all evaluated final configurations reach the 10-epoch cap. The final model uses all applicable training data, seed +1000, no validation set, and no early stopping. Main tables report five-run means and standard deviations rather than final-model point estimates.

Model sizes, software, and hardware environment. The encoders have approximately 178M (mBERT), 278M (XLM-RoBERTa-base), 184M (DeBERTa-v3-base), and 435M (DeBERTa-v3-large) parameters; the SFT generator Qwen3-8B has about 8B parameters, of which QLoRA updates only low-rank adapters. Detector training and inference are performed on a single NVIDIA GeForce RTX 3090 (24 GB); the five-fold training of the backbone comparison, architecture, and ablation variants together with the final model amounts to on the order of tens of GPU-hours, with a single detector run taking from a few minutes to a few hours and SFT about two hours. The software environment is PyTorch 2.9.0 (CUDA 12.8), transformers 4.57.1, and peft 0.18.0. AI texts are generated locally through Ollama (Appendix A.4); blinded LLM evaluation uses API calls made in

May 2026 (Appendix C.3).

B.3 SFT (QLoRA). We fine-tune Qwen3-8B using QLoRA (Detmers et al., 2023), with 4-bit loading and bfloat16 computation. The LoRA adapters (Hu et al., 2022a) use rank 64, alpha 16, dropout 0, and target q/k/v/o_proj plus gate/up/down_proj. The maximum sequence length is 2,048, with gradient checkpointing enabled. Training uses learning rate $5e-5$, three epochs, batch size 1, gradient accumulation 8, and warmup ratio 0.05. After language balancing, a seed-42 split holds out 15% of GenTrain for validation; validation loss selects the checkpoint.

C Complete Experimental Results

C.1 Macro-averaged F1 (supplement to Table 1). RCNN macro-averaged F1, under the same five-run protocol and calculated from each run’s confusion matrix, is 0.6678 ± 0.036 for sarcasm and 0.9298 ± 0.005 for MGTD. Sarcasm macro-F1 is higher than its positive-class F1 (0.6309) because the majority class has a higher F1. These values supplement rather than rerank the full positive-class comparison in §5.1.

C.2 Complete four-quadrant results for the MGTD head (including single-head controls). The reporting convention matches Table 3 (recall for human subsets; false-positive rate for AI subsets).

Model	H×s	H×n	AI×s	AI×n
Plain dual	91.4	91.5	3.6	6.1
Inter.-proj. dual	85.4	85.4	1.8	4.1
RCNN dual	94.8	93.2	5.5	6.6
Plain single	90.4	86.6	4.1	4.2
RCNN single	90.5	92.4	3.6	5.7

Table 7: MGTD-head four-quadrant results (%): recall for human (H) subsets, false-positive rate for AI subsets. s = sarcastic, n = non-sarcastic.

The single-head MGTD controls remain within single-digit percentage points of the dual-head models in every quadrant. No analogous extreme class-conditional error appears on the MGTD side, consistent with §5.3.

No comparable subgroup increase appears on the qwen3.5:9b replication set: AI × non-sarcastic false-positive rates are 0.9%, 1.3%, and 0.5% for plain, interaction projection, and RCNN; AI × sarcastic rates are at most 0.4%; and overall MGTD F1 scores are 0.952, 0.917, and 0.966.

C.3 Blinded-evaluation protocol, inter-judge agreement, and group-level tests (§6). Each sample is sent separately to the three judge APIs with one prompt (May 2026; temperature = 0.3; GPT-5-nano uses its default decoding because it does not support this parameter). The prompt contains only the input and evaluated text, not the ground-truth source. It asks for source inference followed by 1–5 ratings of sarcasm and human-likeness. Base and SFT outputs use the same prompt and are evaluated separately.

Pairwise instance-level agreement is moderate to low. Spearman ρ ranges from 0.20 to 0.38 for sarcasm (all $p < 10^{-6}$) and from 0.03 to 0.23 for human-likeness (one pair is not significant). The conclusion in §6 therefore rests on consistent condition-level directions, not agreement in instance rankings.

Wilcoxon signed-rank tests with Benjamini–Hochberg FDR correction (Benjamini and Hochberg, 1995) show that SFT receives lower sarcasm ratings than the base model in five of six language $\times$ judge groups. The pooled correlations in Table 8 use the three-judge mean as a descriptive aggregate; because inter-judge agreement is limited, the judge-specific results remain necessary. Four unparseable ratings occur across four instances, each of which retains two valid judge ratings; the aggregate uses the mean of available ratings and therefore retains all 600 outputs.

Subset	n	Sarc. ρ	MGTD ρ
Overall	600	+0.136*	−0.212*
Chinese	360	−0.020	−0.218*
English	240	+0.338*	−0.270*

Table 8: Spearman correlations between detector probabilities and LLM-judge ratings (RCNN detector, $n = 600$ pooled over base and SFT). * significant; other entries not significant.

C.4 Paired tests between SFT and human references (supplementary). This test is performed on the Chinese subset of the generation-evaluation set ($n = 180$; GPT-5-nano is missing three instances, leaving 177). Directions are inconsistent: for DeepSeek-V3, SFT is significantly below the reference on human-likeness (-0.161 , $p = 0.0025$); for GPT-5-nano, SFT is significantly above the reference on sarcasm ($+0.186$, $p = 0.0114$); the other four groups are not significant.

Judge / lang.	n	SFT	Base	adj. p
DeepSeek-V3 / en	120	3.292	4.025	<.0001
DeepSeek-V3 / zh	180	3.878	4.044	.0008
Gemini-2.5-FL / en	120	3.750	3.992	.0008
Gemini-2.5-FL / zh	180	3.994	4.028	.2008*
GPT-5-nano / en	119	3.681	4.084	<.0001
GPT-5-nano / zh	180	3.517	3.911	<.0001

Table 9: Group-level tests for blinded sarcasm judgments (Wilcoxon + BH-FDR). * not significant. GPT-5-nano lacks one English base rating (C.3: $n = 119$) and three Chinese reference ratings (affecting C.4, not the C.3 Chinese row); tests use pairwise deletion.

C.5 Plain-detector control. With the plain detector, Chinese sarcasm changes from 0.765 to 0.763 (-0.001), English sarcasm from 0.843 to 0.659 (-0.184), and the MGTD score rises by $+0.760$ in Chinese and $+0.705$ in English. The two detectors therefore agree on the rise in MGTD scores, while the sarcasm direction depends on the detector.

Correlations are $+0.142$ for sarcasm and -0.214 for MGTD overall; -0.037 (not significant) and -0.209 in Chinese; and $+0.342$ and -0.283 in English. Each differs from the RCNN-detector value by at most 0.02. Negative human-authorship-probability versus human-likeness associations are similar across the two detectors, whereas the condition-level sarcasm-score change is architecture-dependent, especially in Chinese.

Within-condition correlations (RCNN detector). For MGTD, the base condition has $\rho = +0.03$ ($n = 300$, not significant; Chinese $+0.10$, English -0.09). The SFT condition has $\rho = -0.12$ ($p = 0.039$; Chinese -0.19 , $p = 0.009$; English -0.12 , not significant). The pooled value of -0.212 is driven mainly by the opposing condition-level changes: SFT raises the detector score by about 0.67 while lowering the mean blinded rating. There is no positive detector–judge alignment on this dimension within conditions.

For sarcasm, within-condition correlations are $+0.05$ for Chinese base (not significant), $+0.15$ for Chinese SFT ($p = 0.048$), $+0.41$ for English base, and $+0.37$ for English SFT (both English values significant). Chinese alignment is therefore at most weakly positive, consistent with its near-zero pooled result.

C.6 Detector scores before and after SFT across the two task–language settings. Scores from

the RCNN detector for base and SFT outputs (0–1):

Lang.	Cond.	Sarc.	MGTD
Chinese	Base → SFT	0.780 → 0.895	0.289 → 0.969
English	Base → SFT	0.854 → 0.730	0.198 → 0.864

Table 10: Detector scores before/after SFT (RCNN detector).

Sarcasm scores move in opposite directions in the Chinese comment-generation and English rewriting settings: +0.115 and −0.124, respectively. GenTrain itself is balanced at 1,214 instances per language; the evaluation settings instead differ in task, prompt format, source data, and sample size, all of which may contribute to the contrast. MGTD scores rise sharply in both settings, so the detector assigns SFT outputs higher probabilities of human authorship.

Mean LLM judgments for sarcasm / human-likeness are 3.77/4.17 for human references, 4.01/4.24 for the base model, and 3.71/4.08 for SFT. These directions conflict with the detector’s MGTD score, as analyzed in §6.

C.7 Conditional test-instance uncertainty: confidence intervals, paired tests, and subset F1. Conditional on each fixed final model, we quantify test-instance uncertainty with Wilson 95% intervals for quadrant rates and 95% percentile intervals from 10,000 instance-level bootstrap samples for F1. These intervals do not include uncertainty from training, generator sampling, corpus construction, or human–AI pairing. They complement the between-run standard deviations in Table 2 and Table 1. All values use the final model for each architecture and follow the reporting convention of Table 3.




Figure 2: Sarcasm false-positive rates on the AI × non-sarcastic subset for five detectors and two generators. Error bars are Wilson 95% intervals conditional on each fixed model.

The fixed final-model configurations show the same directional dual–single gap with the second generator, although its magnitude differs.

Model	deepseek	qwen
Plain dual	86.3 [84.5, 88.0]	91.3 [89.8, 92.6]
Inter.-proj. dual	40.9 [38.5, 43.4]	29.5 [27.3, 31.8]
RCNN dual	34.1 [31.8, 36.5]	22.1 [20.1, 24.3]
Plain single	19.3 [17.4, 21.4]	7.9 [6.6, 9.3]
RCNN single	19.3 [17.4, 21.4]	10.6 [9.2, 12.3]

Table 11: Wilson 95% intervals for the sarcasm false-positive rate on the AI × non-sarcastic subset (two generators; %).

Paired McNemar tests use predictions on the same AI × non-sarcastic texts from the main-generator set. Discordant pairs are 1,021 / 7 for plain dual vs. plain single, 239 / 15 for RCNN dual vs. RCNN single, 805 / 15 for plain dual vs. RCNN dual, and 203 / 100 for interaction projection vs. RCNN dual (all $p < .001$). The discordant counts and absolute rate differences carry the substantive effect size.

Model	Lang.	n	Dual pos.; rate [CI]	Single pos.; rate [CI]	Δ pp
Plain	en	1,174	987; 84.1 [81.9, 86.1]	242; 20.6 [18.4, 23.0]	+63.5
Plain	zh	339	319; 94.1 [91.1, 96.2]	50; 14.7 [11.4, 18.9]	+79.4
RCNN	en	1,174	445; 37.9 [35.2, 40.7]	247; 21.0 [18.8, 23.5]	+16.9
RCNN	zh	339	71; 20.9 [17.0, 25.6]	45; 13.3 [10.1, 17.3]	+7.7

Table 12: Language-stratified sarcasm false-positive rates on AI texts prompted to be non-sarcastic (%; Wilson 95% intervals). Exact paired McNemar tests give $p < .001$ for both plain comparisons and RCNN English, and $p = 2.4 \times 10^{-5}$ for RCNN Chinese.

Both same-architecture comparisons therefore have the same direction in Chinese and English. Language remains coupled with corpus, genre, and annotation setting, so these results provide stratified replication across the two task–language settings rather than an isolated language effect.

Main-model MGTD F1 is 0.938 [0.931, 0.945] on the full set, 0.961 [0.952, 0.969] in Chinese, and 0.923 [0.912, 0.933] in English. Because the source label is constant within the human subset, positive-class F1 reduces to recall and is not reported separately.

Chinese sarcasm F1 is higher than English sarcasm F1 for the RCNN model (0.819 vs. 0.429). The contrast is confounded by prevalence, genre, annotation procedure, and language: sarcasm prevalence is 63.8% in Chinese but 14.2% in English, and the English human side is dominated by author-labeled intended-sarcasm tweets from iS-

Model	Full set	Chinese	English	Human subset	Human-zh	Human-en
Plain dual-head	0.556	0.746	0.312	0.586	0.710	0.374
Interaction-projection dual-head	0.643	0.781	0.420	0.593	0.739	0.366
RCNN dual-head	0.663 [.646, .680]	0.819 [.802, .836]	0.429 [.399, .458]	0.629 [.605, .653]	0.774 [.748, .798]	0.399 [.355, .440]
Plain single-head	0.661	0.750	0.514	0.612	0.719	0.434
RCNN single-head	0.664	0.758	0.511	0.603	0.709	0.434

Table 13: Per-language and human-subset sarcasm F1 (final models; bootstrap 95% intervals included for RCNN).

arcasmEval whereas the Chinese side uses topic-comment-style ToSarcasm. The positive-class F1 values are therefore not directly comparable as a pure language contrast.

DeBERTa-v3’s SentencePiece vocabulary (Kudo and Richardson, 2018) tokenizes Chinese largely by character, at approximately 1.02 tokens per character with 2.5% out-of-vocabulary symbols. This observation does not separate language from the other dataset differences above.

The separately retrained final RCNN model’s overall sarcasm F1 is 0.663, compared with 0.629 on human texts alone, because the overall score also includes AI texts. The five-run mean in Table 1 is a different estimate. Within the human subset, dual–single differences are below 0.03 and vary in direction. The visible negative-transfer loss is therefore concentrated in the AI-side quadrants, consistent with §5.3.

False-positive rates use a threshold of 0.5. For the plain dual-head model on AI $\times$ non-sarcastic text, the median sarcasm probability is 0.767 with an interquartile range of 0.635–0.839. The distribution therefore shifts across the threshold rather than placing only a few cases near it. Single-head medians are 0.204 for the plain model and 0.069 for RCNN.

AUC values within AI texts are similar across models (0.82–0.88). Joint training thus raises sarcasm scores while largely preserving their internal ordering. This pattern is consistent with a source-related component being added to the sarcasm signal.

D Representative Cases of Detector–LLM Score Misalignment

Cases were manually selected after inspecting the paired table of RCNN-detector scores and ratings from the three LLMs to illustrate large disagreements; no formal ranking or top- k selection was implemented. Detector scores range from 0 to 1, and LLM scores are the three-evaluator mean on a 1–5 scale; long texts are truncated. Chinese examples retain the original model input and provide a

pinyin transliteration and an English translation in brackets.

D.1 Classified as human-authored by the detector but rated as less human-like by the LLMs (MGTD-score misalignment; corresponding to the analysis in §6). In the first case, AI-generated text with obvious chain-of-thought traces is nevertheless assigned human-authorship probability 1.00 by the MGTD head, corresponding to the score disagreement discussed in §6.

D.2 Detector-high / LLM-judge-low sarcasm cases.

D.3 Detector-low / LLM-judge-high sarcasm cases. Together, these selected cases motivate the hypothesis that surface markers such as mock-praise formulas, exclamation marks, and colloquial short sentences contribute to some detector decisions, while semantic reversal contributes to some LLM-judge ratings. Systematic feature-level testing is required to determine whether either pattern generalizes.

The manually selected cases do not estimate how often either pattern occurs. They provide qualitative examples of the near-zero Chinese sarcasm correlation in §6; systematic testing remains future work.

L	S	$P_{\text{hum}}/\text{LLM}$	Text (excerpt)
zh	SFT	1.00 / 2.67	“好，我需要针对三星回怼华为的新闻生成一条讽刺评论。首先……” [hǎo, wǒ xūyào zhēnduì Sānxīng huíduì Huáwéi de xīnwén shēngchéng yī tiáo fěngcì pínglùn. shǒuxiān...] [“Okay, I need to generate a sarcastic comment on the news of Samsung hitting back at Huawei. First...”] (chain-of-thought leakage)
zh	SFT	1.00 / 2.33	“跑分竞赛：参数焦虑：跑分焦虑：水货：哈哈哈” [pǎofēn jìngsài: cānshù jiāolǜ: pǎofēn jiāolǜ: shuǐhuò: hāhāhā] [“Benchmark race: spec anxiety: benchmark anxiety: grey-market goods: hahaha”] (degenerate output)
en	SFT	1.00 / 2.67	“Same as me, I’m not sure” (vacuous short sentence)

Table 14: MGTD-score misalignment cases. $P_{\text{hum}} = P(\text{human})$.

L	S	$P_{\text{sar}}/\text{LLM}$	Text (excerpt)
zh	SFT	1.00 / 2.67	“三星：我才是第一！你看看我这摄像头总分，109+96，112+89，我赢了！” [Sānxīng: wǒ cái shì dì-yī! nǐ kànkān wǒ zhè shèxiàngtóu zǒngfēn, 109+96, 112+89, wǒ yíng le!] [“Samsung: I’m the real number one! Look at my camera total scores, 109+96, 112+89, I win!”] (numbers repeated without sarcastic transformation)
zh	SFT	0.99 / 2.67	“恭喜恭喜，美国牛牛牛，继续加油” [gōngxǐ gōngxǐ, Měiguó niú niú niú, jìxù jiāyóu] [“Congratulations, congratulations, America is awesome awesome awesome, keep it up”] (formulaic mock-praise)
en	SFT	0.93 / 2.67	“2 hours left to tip it over, I’m going to the toilet”

Table 15: Detector-high / LLM-judge-low sarcasm cases. $P_{\text{sar}} = P(\text{sarcastic})$.

L	S	$P_{\text{sar}}/\text{LLM}$	Text (excerpt)
zh	Base	0.14 / 4.00	“特朗普承诺 120 亿补贴像撒大饼，农民收到的却是缩水成硬币的承诺……” [Tèlǎngpǔ chéngnuò 120 yì bǔtiē xiàng sā dàbǐng, nóngmín shōudào de què shì suōshuǐ chéng yìngbì de chéngnuò...] [“Trump promised a 12-billion subsidy as if tossing out big pancakes, yet what farmers receive are promises shrunk into coins...”] (metaphorical irony)
zh	Base	0.24 / 4.00	“特斯拉：价格跳水后又跳水，消费者：钱包跳水后又跳水。” [Tèsīlā: jiàgé tiàoshuǐ hòu yòu tiàoshuǐ, xiāofèizhě: qiānbāo tiàoshuǐ hòu yòu tiàoshuǐ.] [“Tesla: prices dive and then dive again; consumers: wallets dive and then dive again.”] (parallel-structure irony)
en	Base	0.12 / 4.00	“Thank you so much for this *beloved* review, Mandy! I’m *overjoyed* you enjoyed it” (asterisk emphasis)

Table 16: Detector-low / LLM-judge-high sarcasm cases. $P_{\text{sar}} = P(\text{sarcastic})$.

E Preliminary Representation-Level Analysis of the Failure Mechanism

This appendix records the complete protocol, data, and self-checks for the representation-level tests in §5.3.2. The interpretation of conclusions follows §5.3.2.

E.1 Setup. All analyses freeze the plain dual-head model that produces the 86.3% main-text result. We use split-sample evaluation: one half estimates the direction and the other computes metrics. The appendix baseline is therefore 84.8% on the evaluation half rather than 86.3% on the full set.

Implementation checks pass. Manually computed logits match model outputs with maximum error 1.4×10^{-6} , and the erasure-projection residual is about 10^{-6} . LEACE satisfies its defining property: after erasure the fitting-half class-mean difference has relative residual 7×10^{-6} , and an in-sample linear probe is at chance (AUC 0.514). Probe selectivity is 0.986 with true labels and 0.512 with permuted labels.

E.2 Protocol. The linear probe is logistic regression with five-fold cross-validation. A selectivity control retrains it with permuted source labels, and correlation analyses are repeated within text-length tertiles. For alignment, Σ is the Ledoit–Wolf covariance fitted to all pooled Dual-Test representations, with eigenvalues clipped below 10^{-8} ; w denotes a task-head weight differ-

ence or probe direction. Whitened alignment is $\cos(\Sigma^{1/2}w_1, \Sigma^{1/2}w_2)$.

Erasure tests cover one direction, INLP subspaces ($k = 1-16$; Ravfogel et al., 2020), and LEACE (Belrose et al., 2023). Controls use ten sampled random whitened directions and directions estimated from permuted labels. srcORTH-PC1 is the maximum-variance direction after removing its projection onto the estimated linear source direction. Erasing it leaves measured linear source decodability unchanged, although orthogonality to one estimated direction does not imply independence from all multidimensional or nonlinear source information.

E.3 Diagnostic evidence. Source is highly linearly decodable: AUC is 0.986 overall, 0.990 on sarcastic texts, and 0.985 on non-sarcastic texts. The permuted-label control gives 0.512. The single-head model also gives 0.974, so source decodability is not unique to joint training. On AI $\times$ non-sarcastic texts, the MGTD-head AI logit and sarcasm-head logit correlate at $\rho = 0.519$; correlations by text-length tertile are 0.44, 0.56, and 0.63.

The sarcasm and source directions are nearly orthogonal in the original geometry ($\cos \approx -0.03$) but moderately aligned after whitening. The absolute cosine is 0.41 using the probe direction and 0.49 using the MGTD-head direction. These results are compatible with the shortcut hypothesis, but remain correlational.

E.4 Intervention evidence. $\text{gap} = \text{AI} \times \text{non-sarcastic false-positive rate} - \text{human} \times \text{non-sarcastic false-positive rate}$; source AUC is the decoding ability of a probe retrained after erasure.

Erasure scheme	AI $\times$ n FP	H $\times$ n FP	Gap	AUC
None (baseline)	84.8	32.8	52.0	0.982
Single dir. (head / probe / random)	84.8–86.2	—	—	≈ 0.98
INLP ($k = 16$)	89.8	—	—	0.974
LEACE (linear)	68.7	49.5	19.1	0.589
Random	84.9	32.7	—	0.982
whitened $\times 10$				
srcORTH-PC1	66.7	54.1	12.6	0.982

Table 17: Intervention evidence. H $\times$ n = human $\times$ non-sarcastic; AI $\times$ n = AI $\times$ non-sarcastic; FP in %.

Single-direction and low-dimensional INLP erasure produce little change. LEACE reduces measured linear source decodability and the source-specific gap from 52.0 to 19.1; the ten sampled random-direction controls show no comparable

gap change.

The orthogonal high-variance control changes the interpretation. It leaves measured linear source decodability intact but changes the AI-side rate to 66.7 and the source gap to 12.6, an effect similar to LEACE. This shows that predictions are sensitive to at least one high-variance direction after removing its projection onto the estimated source direction; it does not identify an independent or unique causal subspace.

Under LEACE, the unweighted mean false-positive rate across the two non-sarcastic source groups is almost unchanged ($58.8 \rightarrow 59.1$). The intervention therefore reduces the source disparity primarily by shifting errors from AI texts to human texts rather than improving non-sarcastic classification overall.