

Joint Training with Machine-Generated Text Detection Can Distort Sarcasm Detection: Diagnosing Class-Conditional Negative Transfer

Anonymous ACL submission

Abstract

Sarcasm detection requires intent beyond literal wording. LLM generation can expand sarcasm resources, but evaluation must distinguish sarcasm-relevant pragmatic cues from source-related cues. A shared encoder supports both sarcasm detection and machine-generated text detection (MGTD), yet source-predictive cues may distort sarcasm predictions. We harmonize Chinese and English resources and construct DualTest for source $\times$ sarcasm subgroup analysis, with prompt-specified style serving as the sarcasm label for AI-generated texts. On AI-generated texts prompted to be non-sarcastic, the baseline joint procedure’s false-positive rate exceeds its sarcasm-only counterpart by 67.0 percentage points, a pattern we call class-conditional negative transfer. Within AI-generated texts, AUC remains 0.82–0.88, showing preserved ranking, while score distributions shift upward. The higher false-positive rates appear in both languages and persist when the same detectors are evaluated, without retraining, on outputs from another generator. Representation diagnostics link source-related features to the affected sarcasm decisions and show that removing one data-derived high-variance direction changes the subgroup gap. Exploratory analysis finds that predicted human-authorship probabilities do not consistently track LLM-judge ratings of human-likeness obtained without ground-truth source labels. These findings motivate source $\times$ sarcasm subgroup evaluation and external validation before detector probabilities are reused as quality scores or rewards.

1 Introduction

Sarcasm expresses an intended attitude that may differ from an utterance’s literal meaning, so automatic detection depends on implicit contrast, communicative intent, and context. Errors can reverse sentiment predictions and disrupt opinion-mining

or content-moderation systems (Tay et al., 2018; Vidgen et al., 2019). Annotation and model behavior also vary with elicitation method, annotator perspective, domain, style, and language (Oprea and Magdy, 2020; Casola et al., 2024; Jang and Frassinelli, 2024).

LLMs offer a way to expand sarcasm resources, but their usefulness depends on realizing the requested style without turning source-specific regularities into shortcuts. Evaluation should therefore separate source and sarcasm: both human-authored and AI-generated texts may be sarcastic or non-sarcastic.

A shared encoder with task-specific heads can reuse one representation for both predictions, but it also creates a route for interference: MGTD favors features that distinguish human-written from AI-generated text (Guo et al., 2023; Wang et al., 2024), whereas sarcasm detection should reflect intent rather than source. The labels themselves do not conflict; the risk lies in how the tasks organize the shared representation. If source-predictive cues influence the sarcasm head, information useful for MGTD but spurious for sarcasm can produce task conflict and negative transfer (Mueller et al., 2022; Hu et al., 2022b). This motivates testing how joint training changes sarcasm predictions.

Because source-predictive regularities and sarcasm cues vary across languages and corpora (Jang and Frassinelli, 2024; Casola et al., 2024; Wang et al., 2024), we examine the transfer pattern separately in Chinese and English.

DualTest supports source $\times$ sarcasm analysis, with prompt-specified style serving as the AI-text sarcasm label. Against same-configuration sarcasm-only controls, joint training sharply increases false positives on AI-generated texts prompted to be non-sarcastic. We call this class-conditional negative transfer because the degradation is concentrated in one source $\times$ sarcasm sub-

group rather than distributed uniformly across DualTest. The effect appears in both languages and persists when the same detectors evaluate outputs from another generator without retraining. Within AI-generated texts, sarcasm AUC remains 0.82–0.88, showing preserved ranking, while score distributions reveal an upward shift.

To connect decodable information with model behavior beyond probe accuracy (Hewitt and Liang, 2019; Elazar et al., 2021), we combine probes, task-head correlations, and representation interventions. Source-related features align with the affected sarcasm decisions, and removing one data-derived high-variance direction changes the subgroup gap, making its content and causal role concrete targets for further analysis.

Detector scores can also enter generation rewards: ViSP includes a DIP-derived sarcasm score in its PPO reward (Wang et al., 2026; Wen et al., 2023). This motivates comparing detector sarcasm probabilities with LLM-judge sarcasm ratings. Separately, because predicted human-authorship probability does not measure perceived human-likeness, we compare those two signals as well, using ratings obtained without ground-truth source labels.

Contributions. We construct DualTest from Chinese and English resources and use same-configuration controls to identify class-conditional negative transfer concentrated on AI-generated texts prompted to be non-sarcastic. The effect persists across both languages and when the same detectors evaluate another generator’s outputs without retraining, while representation diagnostics link source-related features to the affected decisions and identify a direction for further causal analysis. Comparing detector probabilities with LLM-judge ratings further illustrates the risk of treating them as measures of sarcasm quality or human-likeness.

2 Related Work

2.1 Sarcasm Detection, Synthetic Data, and Multilingual Resources

Feature-engineered methods (Riloff et al., 2013), neural sequence models (Ghosh and Veale, 2016; Tay et al., 2018), and pretrained language models (Khatri and P, 2020; Du et al., 2022) have each advanced sarcasm detection. Performance nevertheless varies substantially across datasets,

which differ in label definitions, domains, languages, and modalities. Label schemes may reflect intended or annotator-perceived sarcasm (Oprea and Magdy, 2020; Abu Farha et al., 2022; Casola et al., 2024), while multimodal resources may use modality-specific annotations (Yue et al., 2024).

Recent studies evaluate LLM-generated sarcasm through predictive performance, diversity, linguistic properties, and human judgments (Majer et al., 2026), and associate synthetic-data characteristics with shortcut learning and poor transfer to human-authored text (Mahato et al., 2026). These findings motivate testing whether MGTD-derived source cues interfere with sarcasm detection in a shared encoder.

2.2 Machine-Generated Text Detection and Authorship Verification

MGTD includes likelihood-based zero-shot methods (Mitchell et al., 2023) and discriminative models trained on paired human–AI corpora (HC3; Guo et al., 2023). M4 documents generalization gaps across generators, domains, and languages (Wang et al., 2024), whereas RAID stress-tests unseen generators, decoding variations, and attacks (Dugan et al., 2024). These findings are consistent with reliance on generator- or domain-specific source cues, which may also shape a shared representation used for sarcasm detection. Bevendorff et al. (2025) frame realistic LLM detection as open-set authorship verification rather than closed-set attribution; here, binary human-versus-AI classification serves as a controlled auxiliary task.

To the best of our knowledge, the interaction between sarcasm detection and MGTD under shared-encoder joint training has not been directly examined. Because source is the MGTD target but should not determine sarcasm, we test whether the shared representation preserves that distinction.

2.3 Multi-Task Negative Transfer, Shortcut Learning, and Shared-Encoder Mitigation

In multi-task learning, transfer depends on task relatedness and data-distribution alignment (Zhang et al., 2023; Wu et al., 2020), while gradient conflict can contribute to negative transfer (Yu et al., 2020; Liu et al., 2021). Matched single- and multi-task comparisons show similar failures and suggest that models may infer task identity from input distributions (Mueller et al., 2022); other work analyzes functional specialization in multi-

task Transformers (He and Choi, 2021; Li et al., 2023). A complementary shortcut-learning account is that models exploit training-predictive features that fail under distribution shift (Geirhos et al., 2020; Ye et al., 2026). Shared representations can strengthen such correlations (Hu et al., 2022b; Chai et al., 2025), while additional datasets can introduce them and reduce worst-group performance (Compton et al., 2023).

Mitigation approaches either isolate task-specific parameters (Pfeiffer et al., 2021, 2020; Karimi Mahabadi et al., 2021; Stickland and Murray, 2019) or modify information flow within a shared graph (Wu et al., 2019; Pilault et al., 2021). Our comparison of a plain shared encoder with two intermediate-representation variants follows the second direction and examines subgroup variation in negative transfer.

2.4 Detector Outputs as Evaluation and Reward Signals

Reusing detector outputs as evaluation or reward signals creates a broader risk. LLM judges increasingly evaluate generated text (Zheng et al., 2023), while optimization against model-based rewards can improve a proxy but degrade a gold-standard measure (Gao et al., 2023). High classification accuracy does not establish that a detector’s continuous output measures style strength or generation quality. For example, ViSP (Wang et al., 2026) includes a DIP-derived sarcasm score (Wen et al., 2023) in its PPO reward (Schulman et al., 2017). We therefore compare detector sarcasm probabilities with separately obtained sarcasm ratings. In parallel, we compare predicted human-authorship probabilities with ratings of perceived human-likeness, two related but non-equivalent constructs.

3 Dataset Construction

Testing cross-task interference requires distinct source and sarcasm supervision, with both attributes labeled for evaluation. Here “source” indicates whether a text is human-authored or AI-generated, whereas “corpus” refers to the original dataset from which it was drawn. We combine the English iSarcasmEval data (Oprea and Magdy, 2020; Abu Farha et al., 2022), CSC (Jang and Frassinelli, 2024; Jang et al., 2023), ToSarcasm (Liang et al., 2022), News Headlines v2 (Misra and Arora, 2023; Misra and Grover, 2021), and

the Open Chinese Internet Sarcasm Corpus (Zhu, 2022). Because none provides both a source label and a sarcasm label, we map the original annotations to a common binary sarcasm label space and construct 44,862 instances (approximately 87% English). Four resources then serve distinct roles: SarcTrain for sarcasm supervision, AuthTrain for source supervision, GenTrain for generation fine-tuning, and DualTest for source $\times$ sarcasm evaluation (Appendix A).

SarcTrain. SarcTrain contains 10,000 human-authored texts for sarcasm classification (2,500 per language $\times$ label bucket). Three buckets are downsampled, and the Chinese sarcastic bucket is expanded from 2,396 to 2,500 by sampling 104 records with replacement. We then partition the dataset into five folds. Because oversampling precedes fold assignment, some duplicated records occur in different folds, as quantified in Appendix A.

AuthTrain. AuthTrain contains 19,995 texts for source classification. Seed-level quality control excludes one Chinese SarcTrain instance, leaving 9,999 approximately language-balanced human seeds (5,000 English; 4,999 Chinese). The locally deployed model identified by the Ollama tag deepseek-r1:8b (Guo et al., 2025; Ollama, 2025) generates the AI counterparts using four language $\times$ style templates, with each seed’s prompt style selected from its sarcasm label. After generation and quality control, 9,996 counterparts are retained, leaving three unmatched seeds and 99.97% pairing coverage. Source and prompt-specified style are thus distinct, approximately balanced factors for source $\times$ sarcasm subgroup analysis (Appendix A).

GenTrain. GenTrain contains 2,428 texts for the supervised fine-tuning (SFT) used in §6. From 5,257 eligible records labeled as sarcastic in iSarcasmEval and ToSarcasm, we sample 2,103 (40%; 1,214 English and 889 Chinese). We then sample 325 additional Chinese records with replacement, yielding 1,214 instances per language. By prompt type, the set contains 374 rewriting, 839 response, and 1,215 comment instances (Appendix A).

DualTest. DualTest contains 4,664 texts for source $\times$ sarcasm evaluation. Its human-authored subset combines the official test sets of iSarcasmEval (1,400 English texts) and ToSarcasm (972 Chinese texts). After generation and quality control, 2,292 of the 2,372 human test texts retain AI-

generated counterparts (96.6% coverage). Human-authored texts retain their original dataset annotations, whereas the sarcasm label for AI-generated texts records prompt-specified style rather than independent human annotation of perceived sarcasm.

The four groups contain 823 human-sarcastic and 1,549 human-non-sarcastic texts, plus 779 AI-generated texts prompted to be sarcastic and 1,513 prompted to be non-sarcastic. Record-level checks find no overlap of SarcTrain or AuthTrain with either side of DualTest. Appendix A reports these checks and the language $\times$ source $\times$ sarcasm distribution.

4 Joint and Single-Task Models

4.1 Tasks and Architecture

Given a Chinese or English input text x , the model predicts a binary sarcasm label and a binary source label, with all four label combinations treated as valid. The tasks share a single DeBERTa-v3-large encoder (He et al., 2023), allowing us to test how joint MGTD–sarcasm training changes the representation read by the sarcasm head. We compare three representation configurations and examine how subgroup errors vary across them. Figure 1 shows the RCNN configuration, whose intermediate module combines recurrent context with convolutional pooling for text classification (Lai et al., 2015), a design previously evaluated for multilingual sarcasm detection (Yang and Ikeda, 2024).

The RCNN configuration applies a single-layer BiLSTM followed by parallel one-dimensional convolutions and global max pooling. With hidden size 256 per direction, the BiLSTM yields a 512-dimensional vector at each token. We concatenate it with the 1,024-dimensional encoder output and apply kernels of widths 3, 4, and 5, each with 128 output channels. Global max pooling yields a 384-dimensional vector z , which passes through dropout to the two heads.

The plain configuration, used as the shared-encoder baseline, sends [CLS] directly to both heads. The second configuration uses a module termed task-branch interaction projection, which forms learned 256-dimensional sarcasm and authorship (MGTD) branch vectors h_s and h_a , and updates them as $h'_s = h_s + P_{a \rightarrow s}(h_a)$ and $h'_a = h_a + P_{s \rightarrow a}(h_s)$. Subscripts identify the branches, arrows give the mapping direction, and primes mark the updated vectors. Because each branch




Figure 1: RCNN dual-head configuration: a shared DeBERTa-v3-large encoder, an RCNN intermediate module (dashed), and two task-specific heads. The comparisons in §5 remove this component in the plain configuration or replace it with two learned residual projections between the task branches, while retaining the same encoder backbone and task definitions (Appendix B).

supplies a single key–value vector, all attention weights equal one, so the updates are residual cross-branch affine projections in which each P composes the corresponding value and output projections. We refer to the resulting full model as the interaction-projection configuration. All three configurations retain the same encoder and task definitions (Appendix B).

4.2 Training Protocol

We distinguish representation configuration (plain, interaction-projection, or RCNN) from training procedure (joint, sarcasm-only, or MGTD-only): joint models have two task heads, whereas task-specific controls have one active head. We compare joint and sarcasm-only procedures in the plain and RCNN configurations. Within each pair, they share the encoder, representation configuration, sarcasm data, folds, and sarcasm-specific hyperparameters. The joint procedure also receives AuthTrain batches, updates shared parameters from both tasks, follows a mixed-task schedule, and

uses joint early stopping. Compute exposure differs: the control traverses the full sarcasm fold each epoch, whereas the joint procedure samples sarcasm batches within the mixed schedule. The resulting contrast characterizes the complete joint-training procedure.

Training uses task-homogeneous batches in a 2:1 MGTD:sarcasm ratio, approximating the AuthTrain:SarcTrain size ratio. Let $z_s(x)$ and $z_a(x)$ be the representations read by the sarcasm and MGTD heads, and let W_s and W_a be their classifiers; subscripts s and a denote sarcasm and authorship (MGTD). The two representations coincide in the plain and RCNN configurations; in the interaction-projection configuration, they are the task-specific h'_s and h'_a defined above. For active task t , B_t is a task-homogeneous batch of input-label pairs (x, y) , and only its task loss is computed:

$$\mathcal{L}(B_t) = \begin{cases} \frac{1}{|B_t|} \sum_{(x,y) \in B_t} \text{CE}_w(W_s z_s(x), y), & B_t \subset \text{SarcTrain}, \\ \frac{1}{|B_t|} \sum_{(x,y) \in B_t} \text{CE}(W_a z_a(x), y), & B_t \subset \text{AuthTrain}, \end{cases} \quad (1)$$

where CE_w is sarcasm cross-entropy with class-weight vector w (weights are provided in [Appendix B](#)). For the class-balanced SarcTrain dataset, the weighted loss assigns a greater penalty to errors on sarcastic examples. Cross-entropy is computed from logits, with softmax applied only for probability-based evaluation, and each task-homogeneous batch updates the active task head and the shared parameters. The encoder uses a lower learning rate than newly initialized modules. The sarcasm-head rate is fixed across configurations, and [Appendix B](#) reports the MGTD-head rates. Unless otherwise stated, binary predictions use a fixed probability threshold of 0.5 rather than a threshold tuned on DualTest.

For dual-task models, checkpoint selection accounts for both tasks. At epoch e , the early-stopping signal is the equally weighted sum of the two validation losses, each normalized by its first-epoch value:

$$\mathcal{L}_{\text{stop}}^{(e)} = \frac{1}{2} \frac{\mathcal{L}_s^{(e)}}{\mathcal{L}_s^{(1)}} + \frac{1}{2} \frac{\mathcal{L}_a^{(e)}}{\mathcal{L}_a^{(1)}}, \quad (2)$$

where $\mathcal{L}_s^{(e)}$ and $\mathcal{L}_a^{(e)}$ are the epoch- e validation losses for the sarcasm and MGTD tasks on their respective validation sets. Normalization reduces the influence of differences in task-specific loss

scales and learning dynamics. Single-task controls use their own validation loss with the same five-epoch patience ([Appendix B](#)).

We conduct five runs, each using one SarcTrain fold for validation and the remaining four for training, while AuthTrain uses a fixed 90/10 split. Each run reloads the pretrained encoder and randomly initializes new modules and heads. Configurations use the same fold assignments and random-seed values, although parameter initialization may differ across architectures.

DualTest is not included in any training or validation split for either task. Five-run standard deviations reflect joint variation in folds and initialization, which are not separately identified. Wilson intervals are conditional on the predictions of each separately trained final model and quantify uncertainty from the test instances, not uncertainty from training or text generation ([Table 3](#); [Appendix C.8](#)).

For each configuration used in the subgroup analysis, a final model is trained on all applicable data with a fixed epoch count determined from the five validation runs, an independent seed, and no early stopping. DualTest is not used to train these models or determine their training duration, and the final models are not included in five-run aggregates. The exploratory analysis in [§6](#) uses the final RCNN dual-head model ([Appendix B](#)).

5 Experiments and Diagnosis

After training, we use DualTest to compare encoder backbones and representation configurations, evaluate joint- versus single-task procedures, and analyze source $\times$ sarcasm subgroups. We report five-run aggregate comparisons and use separately trained final models for paired subgroup analyses and Wilson intervals.

5.1 Backbone Comparison

We compare dual-head mBERT ([Devlin et al., 2019](#)), XLM-RoBERTa-base ([Conneau et al., 2020](#)), DeBERTa-v3-base, and DeBERTa-v3-large ([He et al., 2023](#)) under one evaluation protocol. For DeBERTa-v3-large, we additionally evaluate the interaction-projection and RCNN representation configurations. On the full DualTest set, we report positive-class F1 for the sarcastic class in sarcasm detection and the human-authored class in MGTD.

[Table 1](#) shows higher mean F1 for the DeBERTa

Model	Sarc. F1	MGTD F1
mBERT dual-head	0.4945 $\pm$.036	0.8743 $\pm$.006
XLm-RoBERTa-base	0.4603 $\pm$.057	0.8321 $\pm$.025
dual-head		
DeBERTa-v3-base dual-head	0.5646 $\pm$.056	0.8952 $\pm$.003
DeBERTa-v3-large plain dual-head	0.6126 $\pm$.031	0.9221 $\pm$.010
+ interaction-projection	0.6076 $\pm$.020	0.9126 $\pm$.007
+ RCNN	0.6309 $\pm$.020	0.9301 $\pm$.005

Table 1: Descriptive dual-head results across backbones and representation configurations (DualTest; five-run mean $\pm$ standard deviation; bold indicates the highest mean).

Config.	Task	Joint	Single-task	Δ
Plain	Sarc.	0.6126 $\pm$.031	0.6734 $\pm$.010	-6.08
Plain	MGTD	0.9221 $\pm$.010	0.9054 $\pm$.027	+1.67
RCNN	Sarc.	0.6309 $\pm$.020	0.6751 $\pm$.022	-4.43
RCNN	MGTD	0.9301 $\pm$.005	0.9253 $\pm$.006	+0.49

Table 2: Joint vs. corresponding single-task procedures (mean $\pm$ standard deviation over five fold-based runs; Δ = joint - single-task, in percentage points).

configurations than for mBERT and XLM-R. The RCNN configuration achieves the highest mean scores in this comparison, reaching 0.6309 sarcasm F1 and 0.9301 MGTD F1, while the two added representation configurations show smaller standard deviations across the five runs than the plain large model. Because backbone family, model scale, and representation configuration vary simultaneously, the comparison does not isolate the effect of any one design choice. For the separately trained final RCNN model, sarcasm F1 is 0.819 in Chinese and 0.429 in English. These language-stratified values are descriptive because prevalence, genre, annotation scheme, and language vary together. Appendix C provides macro-averaged, human-subset, and language-stratified results.

5.2 Joint versus Single-Task Procedures

Each row compares a joint model with the corresponding single-task control for the task shown. For the sarcasm rows, the paired joint and sarcasm-only models share the encoder, representation configuration, sarcasm data, fold assignments, and sarcasm-specific hyperparameters, although their task schedules and sarcasm exposure differ as described in §4.2. Defining negative transfer as lower performance than the corresponding single-task procedure, we find that joint training low-

ers sarcasm F1 in all ten comparisons using the same fold assignments—five per configuration—with differences from -0.8 to -8.6 pp. Mean sarcasm F1 decreases by 6.08 pp for plain and 4.43 pp for RCNN, whereas mean MGTD F1 changes by only +1.67 and +0.49 pp. Sarcasm detection is the only task for which the joint model underperforms its corresponding single-task model in all ten comparisons, and the joint-versus-sarcasm-only gap is smaller for RCNN than for plain. We next examine which source $\times$ sarcasm subgroups account for the aggregate degradation.

5.3 Source $\times$ Sarcasm Subgroup Analysis

Because the aggregate F1 decrease may be uneven across source $\times$ sarcasm groups, we examine the four DualTest subgroups separately. Human-authored texts retain dataset sarcasm annotations, whereas AI-generated texts use prompt-specified style labels. Because each group contains only one reference class, we report recall for sarcastic texts and false-positive rate for non-sarcastic texts.

5.3.1 Three Configurations and Sarcasm-Only Controls

We use the three joint configurations to characterize how subgroup errors vary across representations and the plain and RCNN sarcasm-only models as same-configuration controls. Within each matched pair, both models use the same all-human sarcasm data, while only the joint model is additionally trained on MGTD, so the contrast captures the prediction change between the two training procedures. Table 3 reports sarcasm predictions, and Appendix C provides the corresponding MGTD results.

On AI-generated texts prompted to be non-sarcastic, false-positive rates are 86.3% for the plain joint model, 40.9% for the joint interaction-projection configuration, and 34.1% for the joint RCNN model. The matched plain and RCNN sarcasm-only controls both yield 19.3%, and the latter two joint configurations remain more than 45 pp below the plain joint rate.

On the same texts, 1,021 false positives occur only under the plain joint model, versus 7 only under its matched sarcasm-only control ($p < .001$). The joint RCNN model also has a higher false-positive rate than its matched control, with a smaller gap. The plain and RCNN pairs provide same-configuration contrasts, while the interaction-projection configuration provides a

Model	H sarc. recall	H non-sarc. FPR	AI sarc. recall	AI non-sarc. FPR
Plain dual-head	67.1%	32.9%	98.2%	86.3% [84.5, 88.0]
Interaction-projection dual-head	71.9%	37.6%	94.4%	40.9% [38.5, 43.4]
RCNN dual-head	80.8%	40.5%	89.3%	34.1% [31.8, 36.5]
Plain single-head (sarcasm only)	67.8%	28.7%	76.6%	19.3% [17.4, 21.4]
RCNN single-head (sarcasm only)	66.2%	28.5%	79.1%	19.3% [17.4, 21.4]

Table 3: Sarcasm predictions across four source $\times$ sarcasm groups. Human-authored texts use dataset sarcasm annotations; AI-generated texts use prompt-specified style labels. All dual-head rows use the joint procedure; single-head rows are sarcasm-only controls. Bold marks the focal AI non-sarcastic subgroup. Brackets show Wilson 95% intervals for the separately trained final models; full intervals and paired tests are in [Appendix C.8](#).

complementary comparison in which architecture and optimization also differ.

On the same subgroup, the MGTD head’s human-authorship false-positive rate is 6.1% for plain and 6.6% for RCNN. Within AI-generated texts, sarcasm AUC against the prompt-specified style labels remains 0.82–0.88 across models, so the relative ranking of sarcastic and non-sarcastic texts is largely preserved. The joint procedure instead shifts sarcasm probabilities upward, causing many texts prompted to be non-sarcastic to cross the fixed 0.5 threshold. The subgroup failure therefore consists of an upward shift in sarcasm scores for AI-generated texts while their within-source ranking remains largely preserved.

The matched plain comparison makes the concentration especially clear: its non-sarcastic false-positive rate changes by 4.3 pp on human-authored text but by 67.0 pp on AI-generated text.¹ The aggregate F1 decrease is therefore concentrated in a sharply localized source $\times$ sarcasm subgroup rather than distributed uniformly across DualTest.

5.3.2 Representation-Level Diagnostics

We analyze frozen representations from the separately trained final plain joint model underlying the 86.3% false-positive rate in [Table 3](#). Five-fold source probes achieve AUCs of 0.986 and 0.974 on the joint and matched sarcasm-only representations, showing that source decodability alone does not explain the gap. In the target subgroup, the MGTD-head AI logit correlates with the sarcasm logit ($\rho = 0.519$). Although the MGTD and sarcasm head directions—each defined as the difference between its two class-weight vectors—are nearly orthogonal under Euclidean cosine ($\cos \approx -0.03$), their covariance-normalized alignment is 0.49; the corresponding probe-weight–sarcasm

alignment is 0.41. These results are correlational.

We fit interventions on one half of the frozen representations and evaluate them on the other, where the baseline AI/human non-sarcastic false-positive rates are 84.8%/32.8% and source AUC is 0.982. LEACE ([Belrose et al., 2023](#)), a linear concept-erasure method applied without retraining the sarcasm head, reduces linearly decodable source information: a source probe retrained after erasure gives AUC 0.589, while the rates change to 68.7%/49.5%. Their unweighted mean is nearly unchanged (58.8% to 59.1%), indicating error reallocation rather than overall improvement. Simpler source-direction and INLP removals ([Ravfogel et al., 2020](#)) produce no comparable reduction.

Separately, removing srcORTH-PC1—the maximum-variance direction after projecting out the estimated linear source direction—changes the rates to 66.7%/54.1%, while source AUC remains 0.982; random-direction controls produce no comparable change. We therefore treat srcORTH-PC1 as behaviorally relevant but semantically uncharacterized, not as an identified source feature ([Appendix E](#)).

5.3.3 Evaluation on Qwen3.5-Generated Texts without Detector Retraining

Using the same four language $\times$ style prompt templates, we generate a second AI test subset with Qwen3.5-9B (Ollama tag qwen3.5:9b) ([Qwen Team, 2026](#)) and evaluate it with the same detectors, without retraining them. After quality control, the non-sarcastic subsets generated by deepseek-r1:8b and Qwen3.5-9B contain 1,513 and 1,550 outputs. On Qwen3.5 outputs, the plain joint and sarcasm-only false-positive rates are 91.3% and 7.9%, an observed gap of 83.4 pp compared with 67.0 pp on DeepSeek. The joint RCNN and interaction-projection rates are 22.1% and 29.5%. The higher joint-versus-sarcasm-only rate therefore persists on another generator, while

¹Percentage-point differences are computed from unrounded rates and can therefore differ by 0.1 pp from subtracting the rounded values shown in the tables.

absolute rates and gaps differ (Appendix C.8).

6 Exploratory Comparison of Detector Scores with LLM-Judge Ratings

To test whether detector probabilities agree with independent ratings of generated text, we compare predicted human-authorship and sarcasm probabilities from RCNN with LLM-judge ratings of sarcasm and human-likeness from DeepSeek-V3 (DeepSeek-AI, 2024), Gemini-2.5-Flash-Lite (Google DeepMind, 2025), and GPT-5-nano (OpenAI, 2025) obtained without ground-truth source labels. We evaluate 600 paired outputs—one from base Qwen3-8B (Yang et al., 2025) and one from its QLoRA-adapted version fine-tuned on GenTrain (Dettmers et al., 2023)—for 300 prompts. Prompts comprise 180 Chinese topic-comment and 120 English rewriting tasks, so language and task are not independently varied.

In Chinese topic-comment generation, mean predicted sarcasm probability rises from 0.780 to 0.895 while the mean 1–5 LLM-judge rating falls from 3.99 to 3.80. Spearman ρ between probability and the mean available rating is -0.020 in Chinese and $+0.338$ in English. Predicted human-authorship probability rises in both settings while rated human-likeness declines, yielding a pooled ρ of -0.212 . However, the within-condition values of $+0.03$ for base outputs and -0.12 for SFT outputs show that opposing condition means dominate the pooled association.

The predicted-sarcasm change also differs between settings, increasing by 0.115 for Chinese topic-comment generation but decreasing by 0.124 for English rewriting. Across the six comparisons formed by two settings and three judges, fine-tuned outputs receive lower mean sarcasm ratings in every case, and five differences remain significant after BH-FDR correction. Classification probabilities therefore do not consistently reproduce the ratings’ text-level rankings or condition-level directions.

Because pairwise judge rank correlations are low to moderate, we report judge-specific results and use the ratings for comparison rather than as a gold standard. Replacing the RCNN detector with the plain detector changes the corresponding correlations by at most 0.02. The resulting disagreements motivate external validation before detector probabilities are reused as quality scores or rewards. Appendix C provides the protocol and tests,

and Appendix D presents selected cases.

7 Conclusion

We identify class-conditional negative transfer from joint MGTD–sarcasm training: on AI-generated texts prompted to be non-sarcastic, the baseline joint procedure’s FPR is 67.0 pp above its sarcasm-only counterpart; score distributions shift upward, while AUC within AI-generated texts remains 0.82–0.88. The higher joint-training FPR recurs in both languages and on Qwen3.5 outputs, with the RCNN and interaction-projection configurations showing lower observed false-positive rates on that subgroup. Representation diagnostics link source-related structure to the affected decisions and identify one direction in representation space for further causal analysis, while the divergence between predicted human-authorship probabilities and human-likeness ratings motivates external validation before detector probabilities are used as quality scores or rewards.

Limitations

DualTest combines human sarcasm annotations with prompt-specified style labels for AI-generated texts. Accordingly, false positives for AI-generated texts are defined relative to the style specified in the prompt rather than independent human annotation of perceived sarcasm.

The core joint-versus-sarcasm-only comparisons evaluate the complete joint-training procedure. Within each pair, the architecture, sarcasm data, fold assignments, sarcasm-head learning rate, and class weights are held fixed, while MGTD batches, mixed-task scheduling, shared-parameter updates, and joint early stopping are introduced together. Compute- and update-matched controls could help disentangle these components. Comparisons across architectures do not isolate architecture effects because model capacity and optimization also differ. The cost-sensitive sarcasm loss may also affect absolute recall and false-positive rates.

Because oversampling precedes fold assignment, some duplicated SarcTrain records occur in different folds and may appear on both sides of a training–validation split, introducing dependence into checkpoint selection, although these records do not occur in DualTest. Generation failures from deepseek-r1:8b leave 80 human test texts with-

out retained AI counterparts.

The representation analyses on one separately trained final model connect source-related structure to behavior but do not isolate a unique causal feature or characterize variation across training seeds. The comparison between detector probabilities and LLM-judge ratings does not include human ratings, and language is not varied independently of generation task. Future work should test these findings using independent style annotations, matched training budgets, human ratings, and evaluation across a broader range of generators.

Ethical considerations

All human-authored texts were obtained from publicly released research datasets: iSarcasmEval, CSC, ToSarcasm, News Headlines, and the Open Chinese Internet Sarcasm Corpus. The iSarcasmEval and CSC repositories display MIT licenses, whereas we do not redistribute human-authored text from sources without an identified redistribution license (Appendix A). We collected no new user data or recruited participants, and we did not perform an additional audit of the released corpora for personally identifiable information or offensive content.

All AI-generated texts were produced locally using publicly released open-weight models (deepseek-r1:8b, Qwen3.5-9B, and the base and QLoRA-adapted Qwen3-8B). They were used only for training and evaluation and were labeled as AI-generated. For automated evaluation, evaluation inputs, generated outputs, and rating instructions were transmitted to the DeepSeek-V3, Gemini-2.5-Flash-Lite, and GPT-5-nano APIs. Dataset metadata and ground-truth source labels were not transmitted.

Given the observed subgroup error and the absence of validation for individual-text authorship attribution or high-stakes deployment, detector outputs should not be used as evidence of authorship or misconduct without human review and domain-specific validation. They should not be deployed for content moderation or automated sanctions without evaluating out-of-distribution behavior. The error concentrated on AI-generated texts prompted to be non-sarcastic is an example of systematic subgroup error affecting a particular type of generated content.

AI-assisted tools supported language editing, structural revision, and formatting, but did not gen-

erate the reported experimental outputs or statistics. The authors determined the study design, analyses, and conclusions and take full responsibility for the work.

References

- Ibrahim Abu Farha, Silviu Vlad Oprea, Steven Wilson, and Walid Magdy. 2022. [SemEval-2022 task 6: iSarcasmEval, intended sarcasm detection in English and Arabic](#). In *Proceedings of the 16th International Workshop on Semantic Evaluation (SemEval-2022)*, pages 802–814.
- Nora Belrose, David Schneider-Joseph, Shauli Ravfogel, Ryan Cotterell, Edward Raff, and Stella Biderman. 2023. [LEACE: Perfect linear concept erasure in closed form](#). In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Yoav Benjamini and Yosef Hochberg. 1995. [Controlling the false discovery rate: A practical and powerful approach to multiple testing](#). *Journal of the Royal Statistical Society: Series B (Methodological)*, 57(1):289–300.
- Janek Bevendorff, Matti Wiegmann, Emmelie Richter, Martin Potthast, and Benno Stein. 2025. [The two paradigms of LLM detection: Authorship attribution vs. authorship verification](#). In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 3762–3787.
- Silvia Casola, Simona Frenda, Soda Marem Lo, Erhan Sezerer, Antonio Uva, Valerio Basile, Cristina Bosco, Alessandro Pedrani, Chiara Rubagotti, Viviana Patti, and Davide Bernardi. 2024. [MultiPICO: Multilingual perspectivist irony corpus](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 16008–16021.
- Junyi Chai, Shenyu Lu, and Xiaoqian Wang. 2025. [Identifying and mitigating spurious correlation in multi-task learning](#). In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 25698–25707.
- Rhys Compton, Lily Zhang, Aahlad Puli, and Rajesh Ranganath. 2023. [When more is less: Incorporating additional datasets can hurt performance by introducing spurious correlations](#). In *Proceedings of the 8th Machine Learning for Healthcare Conference*, volume 219 of *Proceedings of Machine Learning Research*, pages 110–127. PMLR.
- Alexis Conneau, Kartikay Khandelwal, Naman Goyal, Vishrav Chaudhary, Guillaume Wenzek, Francisco Guzmán, Edouard Grave, Myle Ott, Luke Zettlemoyer, and Veselin Stoyanov. 2020. [Unsupervised cross-lingual representation learning at scale](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*.

806	DeepSeek-AI. 2024. DeepSeek-V3 technical report .	862
807	arXiv preprint arXiv:2412.19437 .	863
808	Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and	864
809	Luke Zettlemoyer. 2023. QLoRA: Efficient finetun-	865
810	ing of quantized LLMs . In <i>Advances in Neural In-</i>	866
811	<i>formation Processing Systems (NeurIPS)</i> .	
812	Jacob Devlin, Ming-Wei Chang, Kenton Lee, and	867
813	Kristina Toutanova. 2019. BERT: Pre-training of	868
814	deep bidirectional transformers for language under-	869
815	standing . In <i>Proceedings of the 2019 Conference of</i>	870
816	<i>the North American Chapter of the Association for</i>	871
817	<i>Computational Linguistics: Human Language Tech-</i>	
818	<i>nologies (NAACL-HLT)</i> .	
819	Xiyang Du, Dou Hu, Jin Zhi, Lianxin Jiang, and Xi-	872
820	aofeng Shi. 2022. PALI-NLP at SemEval-2022 task	873
821	6: iSarcasmEval- fine-tuning the pre-trained model	874
822	for detecting intended sarcasm . In <i>Proceedings of</i>	875
823	<i>the 16th International Workshop on Semantic Evalu-</i>	876
824	<i>ation (SemEval-2022)</i> , pages 815–819.	877
825	Liam Dugan, Alyssa Hwang, Filip Trhlik, Andrew	878
826	Zhu, Josh Magnus Ludan, Hainiu Xu, Daphne Ip-	879
827	polito, and Chris Callison-Burch. 2024. RAID: A	880
828	shared benchmark for robust evaluation of machine-	881
829	generated text detectors . In <i>Proceedings of the</i>	
830	<i>62nd Annual Meeting of the Association for Compu-</i>	
831	<i>tational Linguistics (ACL)</i> , pages 12463–12492.	
832	Yanai Elazar, Shauli Ravfogel, Alon Jacovi, and Yoav	882
833	Goldberg. 2021. Amnesic probing: Behavioral ex-	883
834	planation with amnesic counterfactuals . <i>Transac-</i>	884
835	<i>tions of the Association for Computational Linguis-</i>	885
836	<i>tics (TACL)</i> , 9:160–175.	886
837	Leo Gao, John Schulman, and Jacob Hilton. 2023. Scal-	887
838	ing laws for reward model overoptimization . In <i>Pro-</i>	888
839	<i>ceedings of the 40th International Conference on</i>	889
840	<i>Machine Learning (ICML)</i> , pages 10835–10866.	890
841	Robert Geirhos, Jörn-Henrik Jacobsen, Claudio	891
842	Michaelis, Richard Zemel, Wieland Brendel,	892
843	Matthias Bethge, and Felix A. Wichmann. 2020.	
844	Shortcut learning in deep neural networks . <i>Nature</i>	
845	<i>Machine Intelligence</i> , 2:665–673.	
846	Aniruddha Ghosh and Tony Veale. 2016. Fracking	893
847	sarcasm using neural network . In <i>Proceedings of</i>	894
848	<i>the 7th Workshop on Computational Approaches to</i>	895
849	<i>Subjectivity, Sentiment and Social Media Analysis</i>	896
850	<i>(WASSA)</i> .	897
851	Google DeepMind. 2025. Gemini 2.5 flash-lite model	898
852	card . Published September 26, 2025.	899
853	Biyang Guo, Xin Zhang, Ziyuan Wang, Minqi Jiang,	900
854	Jinran Nie, Yuxuan Ding, Jianwei Yue, and Yupeng	901
855	Wu. 2023. How close is ChatGPT to human experts?	902
856	comparison corpus, evaluation, and detection . <i>arXiv</i>	903
857	<i>preprint arXiv:2301.07597</i> .	904
858	Daya Guo, Dejian Yang, Haowei Zhang, et al.	905
859	2025. DeepSeek-R1 incentivizes reasoning in LLMs	906
860	through reinforcement learning . <i>Nature</i> , 645:633–	907
861	638.	908
		909
		910
		911
		912
		913
		914
		915
		916

917	classification. In <i>Proceedings of the AAAI Conference on Artificial Intelligence</i> , volume 29, pages 2267–2273.	971
918		972
919		973
920	Olivier Ledoit and Michael Wolf. 2004. A well-conditioned estimator for large-dimensional covariance matrices. <i>Journal of Multivariate Analysis</i> , 88(2):365–411.	974
921		975
922		
923		
924	Chong Li, Shaonan Wang, Yunhao Zhang, Jiajun Zhang, and Chengqing Zong. 2023. Interpreting and exploiting functional specialization in multi-head attention under multi-task learning. In <i>Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)</i> , pages 16460–16476.	976
925		977
926		978
927		979
928		
929		
930		
931	Bin Liang, Zijie Lin, Bing Qin, and Ruifeng Xu. 2022. Topic-oriented sarcasm detection: New task, new dataset and new method (in Chinese). In <i>Proceedings of the 21st Chinese National Conference on Computational Linguistics (CCL)</i> , pages 557–568, Nanchang, China. Chinese Information Processing Society of China.	980
932		981
933		982
934		983
935		984
936		985
937		
938	Bo Liu, Xingchao Liu, Xiaojie Jin, Peter Stone, and Qiang Liu. 2021. Conflict-averse gradient descent for multi-task learning. In <i>Advances in Neural Information Processing Systems (NeurIPS)</i> .	986
939		987
940		988
941		989
942		990
943	Priyanshu Mahato, Aniket Santosh Mishra, and Kripabandhu Ghosh. 2026. LLMs in sarcasm detection? it’s elementary! (or is it?). In <i>Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 29169–29199.	991
944		992
945		993
946		994
947		995
948	Laura Majer, Ana Barić, Florijan Sandalj, Ivan Unković, Bojan Puvača, and Jan Šnajder. 2026. A multi-aspect evaluation framework for synthetic data: Case study on irony and sarcasm. In <i>The Proceedings for the 15th Workshop on Computational Approaches to Subjectivity, Sentiment Social Media Analysis (WASSA 2026)</i> , pages 305–323.	996
949		997
950		998
951		
952		
953		
954		
955	Rishabh Misra and Prahal Arora. 2023. Sarcasm detection using news headlines dataset. <i>AI Open</i> , 4:13–18.	999
956		1000
957		1001
958	Rishabh Misra and Jigyasa Grover. 2021. <i>Sculpting Data for ML: The first act of Machine Learning</i> . Independently published.	1002
959		1003
960		1004
961	Eric Mitchell, Yoonho Lee, Alexander Khazatsky, Christopher D. Manning, and Chelsea Finn. 2023. DetectGPT: Zero-shot machine-generated text detection using probability curvature. In <i>Proceedings of the 40th International Conference on Machine Learning (ICML)</i> .	1005
962		1006
963		1007
964		1008
965		1009
966		1010
967	David Mueller, Nicholas Andrews, and Mark Dredze. 2022. Do text-to-text multi-task learners suffer from task conflict? In <i>Findings of the Association for Computational Linguistics: EMNLP 2022</i> .	1011
968		1012
969		1013
970		1014
	Ollama. 2025. DeepSeek-R1-0528-Qwen3-8B model card. Ollama deployment tag deepseek-r1:8b; accessed 2026-07-21.	1015
		1016
		1017
	OpenAI. 2025. OpenAI GPT-5 system card. August 2025.	1018
		1019
	Silviu Oprea and Walid Magdy. 2020. iSarcasm: A dataset of intended sarcasm. In <i>Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)</i> .	1020
		1021
	Jonas Pfeiffer, Aishwarya Kamath, Andreas Rücklé, Kyunghyun Cho, and Iryna Gurevych. 2021. AdapterFusion: Non-destructive task composition for transfer learning. In <i>Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics (EACL)</i> .	1022
		1023
		1024
	Jonas Pfeiffer, Ivan Vulić, Iryna Gurevych, and Sebastian Ruder. 2020. MAD-X: An adapter-based framework for multi-task cross-lingual transfer. In <i>Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)</i> .	
	Jonathan Pilault, Amine El hattami, and Christopher Pal. 2021. Conditionally adaptive multi-task learning: Improving transfer learning in NLP using fewer parameters and less data. In <i>International Conference on Learning Representations (ICLR)</i> .	
	Qwen Team. 2026. Qwen3.5-9B model card. Official model weights and configuration; accessed 2026-07-21.	
	Shauli Ravfogel, Yanai Elazar, Hila Gonen, Michael Twiton, and Yoav Goldberg. 2020. Null it out: Guarding protected attributes by iterative nullspace projection. In <i>Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)</i> .	
	Ellen Riloff, Ashequl Qadir, Prafulla Surve, Lalindra De Silva, Nathan Gilbert, and Ruihong Huang. 2013. Sarcasm as contrast between a positive sentiment and negative situation. In <i>Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing (EMNLP)</i> .	
	John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. 2017. Proximal policy optimization algorithms. <i>arXiv preprint arXiv:1707.06347</i> .	
	Asa Cooper Stickland and Iain Murray. 2019. BERT and PALs: Projected attention layers for efficient adaptation in multi-task learning. In <i>Proceedings of the 36th International Conference on Machine Learning (ICML)</i> .	
	Yi Tay, Anh Tuan Luu, Siu Cheung Hui, and Jian Su. 2018. Reasoning with sarcasm by reading in-between. In <i>Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (ACL)</i> .	

- Bertie Vidgen, Alex Harris, Dong Nguyen, Rebekah Tromble, Scott Hale, and Helen Margetts. 2019. [Challenges and frontiers in abusive content detection](#). In *Proceedings of the Third Workshop on Abusive Language Online*, pages 80–93.
- Changli Wang, Fang Yin, Jiafeng Liu, and Rui Wu. 2026. [ViSP: A PPO-enhanced framework for multi-modal sarcasm generation with contrastive learning](#). *Neurocomputing*, 678:133185.
- Yuxia Wang, Jonibek Mansurov, Petar Ivanov, Jinyan Su, Artem Shelmanov, Akim Tsvigun, Chenxi Whitehouse, Osama Mohammed Afzal, Tarek Mahmoud, Toru Sasaki, Thomas Arnold, Alham Fikri Aji, Nizar Habash, Iryna Gurevych, and Preslav Nakov. 2024. [M4: Multi-generator, multi-domain, and multi-lingual black-box machine-generated text detection](#). In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1369–1407.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc V. Le, and Denny Zhou. 2022. [Chain-of-thought prompting elicits reasoning in large language models](#). In *Advances in Neural Information Processing Systems (NeurIPS)*, pages 24824–24837.
- Changsong Wen, Guoli Jia, and Jufeng Yang. 2023. [DIP: Dual incongruity perceiving network for sarcasm detection](#). In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 2540–2550.
- Lianwei Wu, Yuan Rao, Haolin Jin, Ambreen Nazir, and Ling Sun. 2019. [Different absorption from the same sharing: Sifted multi-task learning for fake news detection](#). In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*.
- Sen Wu, Hongyang R. Zhang, and Christopher Ré. 2020. [Understanding and improving information transfer in multi-task learning](#). In *International Conference on Learning Representations (ICLR)*.
- An Yang et al. 2025. [Qwen3 technical report](#). *arXiv preprint arXiv:2505.09388*.
- Linshuo Yang and Daisuke Ikeda. 2024. [The influence of linguistic attribute differences in multilingual datasets on sarcasm detection](#). *International Journal of Service and Knowledge Management*, 8(2):1–15.
- Wenqian Ye, Luyang Jiang, Eric Xie, Guangtao Zheng, Yunsheng Ma, Xu Cao, Dongliang Guo, Daiqing Qi, Zeyu He, Yijun Tian, Christopher W. Porter, Megan Coffee, Zhe Zeng, Sheng Li, Ziran Wang, Ting-Hao (Kenneth) Huang, James Matthew Rehg, Henry Kautz, and Aidong Zhang. 2026. [The clever Hans mirage: A comprehensive survey on spurious correlations in machine learning](#). *Transactions on Machine Learning Research*.
- Tianhe Yu, Saurabh Kumar, Abhishek Gupta, Sergey Levine, Karol Hausman, and Chelsea Finn. 2020. [Gradient surgery for multi-task learning](#). In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Tan Yue, Xuzhao Shi, Rui Mao, Zonghai Hu, and Erik Cambria. 2024. [SarcNet: A multilingual multimodal sarcasm detection dataset](#). In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING)*, pages 14325–14335.
- Zhihan Zhang, Wenhao Yu, Mengxia Yu, Zhichun Guo, and Meng Jiang. 2023. [A survey of multi-task learning in natural language processing: Regarding task relatedness and training methods](#). In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics (EACL)*, pages 943–956.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhonghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. [Judging LLM-as-a-judge with MT-Bench and Chatbot Arena](#). In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Yizhang Zhu. 2022. [Open Chinese internet sarcasm corpus construction: An approach](#). *Frontiers in Computing and Intelligent Systems*, 2(1):7–9.

A Detailed Dataset Construction, Generation, and Quality Checks

Corpus	Lang.	Inst.	Resource use
iSarcasmEval (SemEval-2022 T6)	English	3,467	SarcTrain; GenTrain; DualTest*
CSC	English	6,878	SarcTrain
ToSarcasm	Chinese	3,898	SarcTrain; GenTrain; DualTest*
News Headlines v2	English	28,619	SarcTrain
Open Chinese Internet Sarcasm Corpus	Chinese	2,000	SarcTrain

Table 4: Input corpora and resource use. The instance counts refer to the 44,862-record pooled collection; *DualTest uses separate official test splits excluded from these counts. Genres and annotations are iSarcasmEval (tweets, speaker self-annotation), CSC (conversational responses), ToSarcasm (topic comments, third-party annotation), News Headlines (distant supervision from The Onion / HuffPost), and the Open Chinese Corpus (Weibo, third-party annotation).

A.1 Source corpora, sampling, and release scope. The pooled collection contains 44,862 instances in total (86.9% English and 13.1% Chinese). SarcTrain is balanced across the four language $\times$ label groups by downsampling the two English groups and the Chinese non-sarcastic group, then sampling 104 additional Chinese sarcastic records with replacement. The Chinese sarcastic source pool contains 2,396 records corresponding to 2,394 distinct stored text values; after expansion to 2,500 records, its text-duplication rate is 4.24% (A.3).

The iSarcasmEval and CSC repositories display MIT licenses, whereas we did not identify explicit redistribution licenses in the releases we accessed for ToSarcasm, News Headlines v2, or the Open Chinese Internet Sarcasm Corpus. We therefore do not redistribute human-authored text from the latter sources. The planned release is limited to eligible AI-generated text, labels and metadata, and reconstruction scripts that require users to obtain the human-authored corpora from their original sources.

Lang.	Source	Sarc.	Non-sarc.
English	Human	200	1,200
English	AI	192	1,174
Chinese	Human	623	349
Chinese	AI	587	339

Table 5: Language $\times$ source $\times$ sarcasm distribution of DualTest. Human-authored texts use dataset sarcasm annotations; AI-generated texts use prompt-specified style labels.

A.2 Language $\times$ source $\times$ sarcasm distribution of DualTest.

A.3 Data quality checks. Automated checks confirm language and label balance across the five SarcTrain folds. AuthTrain contains 9,996 complete human-AI pairs and three unmatched human instances. Record-level matching finds that neither SarcTrain nor AuthTrain shares records with either the human or AI side of DualTest.

SarcTrain’s text-duplication rate is 0.34% for English and 2.16% for Chinese. Chinese duplication is concentrated in the sarcastic group (4.24%; non-sarcastic group: 0.08%) and results primarily from the oversampling used to fill that group (A.1). Because oversampling precedes fold assignment, for 86 of the 104 added records, another copy originating from the same source record is assigned to

a different fold. The two counting rules are summarized below.

Counting rule	Groups	Records	Cross-fold
Exact stored text; over-sampled copies included	116	241	98
Case/whitespace normalized; added copies excluded	19	46	16

The first rule underlies the reported duplication rates; the second isolates pre-existing duplicates after the added copies are removed. These relationships are relevant to fold-based checkpoint selection but do not place the affected training or validation records into DualTest. Post-processing detects and removes chain-of-thought-style markers (Wei et al., 2022) and the reasoning prefaces described in A.4.

A.4 AI generation protocol. The AI-generated texts in AuthTrain and the two DualTest versions are produced by locally deployed models through the Ollama API. The model deployed under the Ollama tag `deepseek-r1:8b` generates AuthTrain and the DeepSeek version of DualTest’s AI-generated subset. Qwen3.5-9B, deployed under the tag `qwen3.5:9b`, generates a second AI test subset that is evaluated with the same detectors without retraining. Both models use temperature = 0.8, top- p = 0.95 (Holtzman et al., 2020), and a maximum output length of 500 tokens. Generation instructions use four templates defined by language $\times$ style, with style determined by the seed instance’s sarcasm label. The templates constrain output style and length to 20–80 Chinese characters or 10–50 English words. These constraints are intended to reduce, but do not eliminate, source differences in length, so Appendix E separately repeats the representation correlation analysis within length tertiles. The complete Chinese sarcastic template used in the experiment is reproduced verbatim below:

```

【任务】生成一条讽刺风格的评论
【主题】{seed}
【要求】
- 直接给出简短的评论，不要解释思路或分析
- 风格要讽刺、幽默、夸张
- 表达与主题相关的观点即可，不必追求严谨
- 20–80 字为佳
【评论】

```

Its Hanyu Pinyin transliteration is:

```

[Renwu] Shengcheng yi tiao fengci
fengge de pinglun
[Zhuti] {seed}
[Yaoqiu]

```

- Zhijie geichu jianduan de pinglun;
buyao jieshi silu huo fenxi
- Fengge yao fengci, youmo, kuazhang
- Biaoda yu zhuti xiangguan de
guandian ji ke; bubi zhuiqiu yanjin
- 20--80 zi wei jia
[Pinglun]

Its English translation is:

[Task] Generate a sarcastic comment
[Topic] {seed}
[Requirements]
- Directly provide a short comment;
do not explain your reasoning or
analysis
- The comment should be sarcastic,
humorous, and exaggerated
- It is sufficient to express a
viewpoint related to the topic; the
comment need not be rigorous
- Preferably 20--80 Chinese
characters
[Comment]

The Chinese non-sarcastic template differs only in the style line, which requests a normal, natural style. The English templates use the same field structure and request either a sarcastic, humorous, exaggerated style or a normal, natural style, with a length limit of 10–50 words.

Post-processing removes `<think>...</think>` segments, reasoning prefaces such as “好的” (*hǎo de* ‘okay’), “嗯” (*ǎ* ‘hmm’), “Okay,” “Let me,” and “First,” and Markdown markup. Outputs that still exceed the length limit or exhibit cyclic repetition after repeated generation attempts are discarded. AuthTrain retains 9,996 AI counterparts from 9,999 seeds; DualTest retains 2,292 counterparts from 2,372 seeds.

The Qwen3.5-generated AI test subset has no generation failures, yielding 1,550 AI $\times$ non-sarcastic instances rather than 1,513. Generator-specific rates therefore use the full set of retained outputs for each generator rather than an identical seed intersection. For this run, Ollama thinking mode is disabled (`think = false`) to prevent the reasoning process from exhausting the output budget. The same post-processing removes any inline chain-of-thought from the deepseek-r1:8b outputs.

A.5 SFT training and generation evaluation.

After eligibility filtering, 5,257 records labeled as sarcastic remain for generation-task construction (2,972 English; 2,285 Chinese). After shuffling with seed 42, we sample 40% of these records, yielding 2,103 instances (1,214 English; 889 Chinese); language balancing then samples 325 addi-

tional Chinese records with replacement, producing GenTrain (2,428; 1,214 per language). By prompt type, it comprises 374 rewriting instances, 839 response instances, and 1,215 comment instances. Generation evaluation uses a separately constructed set of 300 prompts that is not part of the GenTrain training-validation split (180 Chinese and 120 English): the Chinese task is to generate a sarcastic comment on a given topic, and the English task is to rewrite a normal sentence as a sarcastic one. Language and generation task are therefore not independently varied. This set is used for detector scoring and all three LLM evaluations in §6.

B Hyperparameters and Training Details

Configuration	Enc.	Inter.	Sarc.	MGTD
Plain dual-head	5×10^{-7}	—	5×10^{-6}	1×10^{-5}
Plain	5×10^{-7}	—	5×10^{-6}	—
sarcasm-only				
Plain MGTD-only	5×10^{-7}	—	—	1×10^{-5}
Interaction-projection dual-head	5×10^{-7}	5×10^{-6}	5×10^{-6}	5×10^{-6}
RCNN dual-head	5×10^{-7}	5×10^{-6}	5×10^{-6}	5×10^{-6}
RCNN	5×10^{-7}	5×10^{-6}	5×10^{-6}	—
sarcasm-only				
RCNN	5×10^{-7}	5×10^{-6}	—	1×10^{-5}
MGTD-only				

Table 6: Discriminative learning rates for the dual-head and task-specific controls. The plain setting is also used in the backbone comparison. A dash denotes an inactive module or task head.

B.1 Discriminative learning rates. The sarcasm-head learning rate is fixed at 5×10^{-6} in every configuration with an active sarcasm head. The MGTD-head rate is 1×10^{-5} for the plain dual-head model and the plain and RCNN MGTD-only controls, compared with 5×10^{-6} for the RCNN and interaction-projection dual-head models. Thus, the core joint-versus-sarcasm-only comparisons retain the same sarcasm-specific rates within each architecture. Because the RCNN joint-versus-MGTD-only comparison also changes the MGTD-head learning rate, it is reported descriptively rather than as an optimization-matched estimate of transfer; cross-configuration differences can likewise reflect optimization settings. Early stopping uses the normalized joint validation loss of Equation 2 for dual-head models and the task’s own validation loss for the task-specific controls, with a patience

of five epochs in both procedures.

B.2 Shared training hyperparameters. Training uses `torch.optim.AdamW` with default PyTorch betas (0.9, 0.999) and $\epsilon = 10^{-8}$, weight decay 0.01, and parameter-group learning rates in B.1. We implement a linear 200-step warmup followed by a constant learning rate using PyTorch’s LambdaLR. Training runs for at most 10 epochs, with an early-stopping patience of 5 and gradient clipping at 1.0. The physical batch size is 4 with gradient accumulation over 4 steps, giving an effective batch size of 16 for the active task. The MGTD:sarcasm sampling ratio is 2:1, and the sarcasm class weights are [1.0, 2.5] for the non-sarcastic and sarcastic classes, respectively. The task ratio and class weights were fixed before DualTest evaluation. For the class-balanced SarcTrain dataset, the class weights implement a cost-sensitive objective that assigns a greater penalty to errors on sarcastic examples. The normalized joint validation loss is $0.5 \times [\text{sarcasm validation loss} / \text{first-epoch value}] + 0.5 \times [\text{MGTD validation loss} / \text{first-epoch value}]$. The single-head controls process the full sarcasm training fold each epoch, whereas the dual-head models draw sarcasm batches under the 2:1 task-mixing schedule.

Detector inputs use a maximum sequence length of 512 tokens. In the RCNN module, a width- k convolution uses padding $\lfloor k/2 \rfloor$ followed by ReLU. Global max pooling covers all convolution positions; no post-convolution padding mask is applied.

All fold-based runs use a seed equal to 42 plus the fold index and share the same fold assignments and random-seed values across architectures and ablations. The same fixed AuthTrain 90/10 split is reused in all five runs; only the SarcTrain validation fold changes. For each configuration, the epoch count used to train its final model is the mode of its five best epoch counts plus one, capped at 10 epochs; the resulting counts range from 3 to 10 across configurations. The final model uses all applicable training data, a seed offset of +1000, no validation set, and no early stopping. Main tables report five-run means and standard deviations rather than point estimates from the final models.

B.3 Model sizes and computational environment. The encoders have approximately 178M (mBERT), 278M (XLM-RoBERTa-base), 184M (DeBERTa-v3-base), and 435M (DeBERTa-v3-large) parameters; the SFT generator Qwen3-8B

has about 8B parameters, of which QLoRA updates only low-rank adapters. Detector training and inference are performed on a single NVIDIA GeForce RTX 3090 (24 GB). Across the five-fold backbone, architecture, and ablation experiments and the training of final models, total compute is on the order of tens of GPU-hours. A single detector run takes from a few minutes to a few hours, and SFT takes about two hours. The software environment is PyTorch 2.9.0 (CUDA 12.8), transformers 4.57.1, and peft 0.18.0. AI-generated texts are produced locally through Ollama (Appendix A.4); LLM evaluation uses API calls made in May 2026 (Appendix C.3).

B.4 SFT (QLoRA). We fine-tune Qwen3-8B using QLoRA (Detrmers et al., 2023), with 4-bit loading and bfloat16 computation. The LoRA configuration (Hu et al., 2022a) uses rank 64, alpha 16, and dropout 0.0 and targets `q_proj`, `k_proj`, `v_proj`, `o_proj`, `gate_proj`, `up_proj`, and `down_proj`. The maximum sequence length is 2,048, with gradient checkpointing enabled. Training uses a learning rate of 5×10^{-5} , three epochs, batch size 1, gradient accumulation 8, and a warmup ratio of 0.05. After language balancing, a seed-42 split holds out 15% of GenTrain for validation; validation loss selects the checkpoint. Because language-balancing oversampling precedes the split, duplicated Chinese records may occur in both partitions.

C Complete Experimental Results

C.1 Macro-averaged F1 (supplement to Table 1). Under the same five-run protocol, the RCNN macro-averaged F1 scores computed from each run’s confusion matrix are 0.6678 ± 0.036 for sarcasm and 0.9298 ± 0.005 for MGTD. Sarcasm macro-F1 is higher than its positive-class F1 (0.6309) because the non-sarcastic majority class has a higher F1. These values supplement the positive-class results in §5.1; because macro-F1 is not reported for every configuration, they are not used for cross-configuration ranking.

C.2 Complete four-quadrant results for the MGTD head (including MGTD-only controls). The reporting convention matches Table 3 (recall for human subsets; false-positive rate for AI subsets).

In every subgroup, each MGTD-only control differs from its corresponding dual-head model by fewer than ten percentage points. No comparably

Model	H×s	H×n	AI×s	AI×n
Plain dual-head	91.4	91.5	3.6	6.1
Interaction-projection dual-head	85.4	85.4	1.8	4.1
RCNN dual-head	94.8	93.2	5.5	6.6
Plain MGTD-only	90.4	86.6	4.1	4.2
RCNN MGTD-only	90.5	92.4	3.6	5.7

Table 7: MGTD-head four-quadrant results (%): recall for human (H) subsets, false-positive rate for AI subsets. s = sarcastic, n = non-sarcastic.

large subgroup error appears on the MGTD side, consistent with §5.3.

When the same detectors evaluate the Qwen3.5-generated AI test subset without retraining, the MGTD head again shows no comparable subgroup increase: AI × non-sarcastic false-positive rates are 0.9%, 1.3%, and 0.5% for the plain, interaction-projection, and RCNN configurations; AI × sarcastic rates are at most 0.4%; and overall MGTD F1 scores are 0.952, 0.917, and 0.965.

C.3 Evaluation protocol, inter-judge rank correlations, and group-level tests (§6). Each sample is evaluated separately by each of the three judge APIs in one API call using the same prompt structure for all model conditions (May 2026; temperature = 0.3; GPT-5-nano uses its default decoding because it does not support this parameter). The prompt contains only the input and evaluated text, not the ground-truth source label. It asks for source inference followed by 1–5 ratings of sarcasm and human-likeness. The ratings are therefore obtained without ground-truth source labels, although the judges are explicitly asked to infer source before assigning the two ratings. Base and SFT outputs are evaluated separately.

Pairwise inter-judge rank correlations are low to moderate, with Spearman ρ ranging from 0.20 to 0.38 for sarcasm (all $p < 10^{-6}$) and from 0.03 to 0.23 for human-likeness (one pair is not significant). We therefore do not treat individual-text rankings as a common gold standard. The condition-level observation rests separately on the consistent rating direction across all six language × judge groups.

Outputs from the SFT model receive lower mean sarcasm ratings in all six language × judge groups, and Wilcoxon signed-rank tests with Benjamini–Hochberg FDR correction (Benjamini and Hochberg, 1995) show that five differences re-

main significant. The pooled correlations in Table 8 use the three-judge mean as a descriptive aggregate, while the limited pairwise rank correlations make judge-specific results necessary. Four instances each have one unparsable rating but retain two valid judge ratings, so averaging the available ratings preserves all 600 outputs in the aggregate.

Subset	n	Sarc. ρ	Human-like ρ
Overall	600	+0.136*	−0.212*
Chinese	360	−0.020	−0.218*
English	240	+0.338*	−0.270*

Table 8: Spearman correlations between detector probabilities and LLM-judge ratings (RCNN detector, $n = 600$ pooled over base and SFT): Sarc. ρ compares $P(\text{sarcastic})$ with sarcasm ratings, and Human-like ρ compares $P(\text{human})$ with human-likeness ratings. * significant; the other entries are not significant.

Judge / lang.	n	Base	SFT	adj. p
DeepSeek-V3 / en	120	4.025	3.292	<.0001
DeepSeek-V3 / zh	180	4.044	3.878	.0008
Gemini-2.5-FL / en	120	3.992	3.750	.0008
Gemini-2.5-FL / zh	180	4.028	3.994	.2008*
GPT-5-nano / en	119	4.084	3.681	<.0001
GPT-5-nano / zh	180	3.911	3.517	<.0001

Table 9: Group-level tests for sarcasm ratings obtained without ground-truth source labels (Wilcoxon + BH-FDR). * not significant. GPT-5-nano lacks one English base rating (C.3: $n = 119$) and three Chinese reference ratings (affecting C.4, not the C.3 Chinese row); tests use pairwise deletion.

C.4 Paired tests between SFT outputs and human-authored reference texts (supplementary). These tests are performed on the Chinese subset of the 300-prompt generation evaluation ($n = 180$; three GPT-5-nano ratings are missing, leaving 177). The comparisons provide judge- and dimension-specific comparison context: for DeepSeek-V3, SFT outputs receive significantly lower human-likeness ratings than the reference texts (−0.161, $p = 0.0025$), whereas for GPT-5-nano, SFT outputs receive significantly higher sarcasm ratings than the reference texts (+0.186, $p = 0.0114$). The remaining four comparisons do not reach significance at the reported threshold. Thus, the observed reference differences vary across judges and dimensions; these tests do not estimate a judge-by-dimension interaction.

C.5 Plain-detector control. With the plain detector, the Chinese sarcasm probability changes from 0.765 to 0.763 (-0.001), the English sarcasm probability changes from 0.843 to 0.659 (-0.184), and predicted human-authorship probability rises by $+0.760$ in Chinese and $+0.705$ in English. The two detectors agree on the rise in predicted human-authorship probability. The Chinese sarcasm-probability change is detector-sensitive, whereas both detectors show a decrease in English.

Overall correlations between detector probabilities and corresponding ratings are $+0.142$ for sarcasm and -0.214 for human-likeness. The corresponding values are -0.037 (not significant) and -0.209 in Chinese, and $+0.342$ and -0.283 in English. Each differs from the RCNN-detector value by at most 0.02. Negative associations between predicted human-authorship probabilities and human-likeness ratings are similar across the two detectors, whereas the condition-level sarcasm-probability change differs between the plain and RCNN detectors, especially in Chinese.

C.6 Within-condition correlations (RCNN detector). For predicted human-authorship probability and human-likeness ratings, the base condition has $\rho = +0.03$ ($n = 300$, not significant; Chinese $+0.10$, English -0.09). The SFT condition has $\rho = -0.12$ ($p = 0.039$; Chinese -0.19 , $p = 0.009$; English -0.12 , not significant). The pooled value of -0.212 is driven mainly by the opposing condition-level changes: SFT raises predicted human-authorship probability by about 0.67 while lowering the mean rating. The pooled negative correlation should therefore not be interpreted as a strong instance-level inverse relationship. Within either condition, predicted human-authorship probability is not significantly positively correlated with rated human-likeness.

For sarcasm, within-condition correlations are $+0.05$ for Chinese base (not significant), $+0.15$ for Chinese SFT ($p = 0.048$), $+0.41$ for English base, and $+0.37$ for English SFT (both English values significant). Chinese alignment is therefore at most weakly positive, consistent with its near-zero pooled result.

C.7 Detector probabilities before and after SFT across the two task-language settings. The RCNN detector assigns the following mean probabilities to base and SFT outputs:

Sarcasm probabilities move in opposite direc-

Lang.	Cond.	$P(\text{sarcastic})$	$P(\text{human})$
Chinese	Base $\rightarrow$ SFT	0.780 $\rightarrow$ 0.895	0.289 $\rightarrow$ 0.969
English	Base $\rightarrow$ SFT	0.854 $\rightarrow$ 0.730	0.198 $\rightarrow$ 0.864

Table 10: Mean detector probabilities before and after SFT (RCNN detector).

tions in the Chinese comment-generation and English rewriting settings ($+0.115$ and -0.124 , respectively), while predicted human-authorship probabilities rise sharply in both. GenTrain itself is balanced at 1,214 instances per language, but the evaluation settings differ in task, prompt format, source data, and sample size, all of which may contribute to the contrast.

The mean LLM-judge ratings for the human-authored reference texts are 3.77 for sarcasm and 4.17 for human-likeness; the corresponding values are 4.01 and 4.24 for the base model and 3.71 and 4.08 for SFT. From the base model to SFT, the mean human-likeness rating decreases even as predicted human-authorship probability increases, as analyzed in §6.

C.8 Uncertainty from test instances: confidence intervals, paired tests, and subset F1. For each separately trained final model, we quantify uncertainty from the test instances with Wilson 95% intervals for quadrant rates and 95% percentile intervals from 10,000 instance-level bootstrap samples for F1. The instance-level bootstrap resamples text instances independently and does not preserve human-AI seed pairs. These intervals do not include uncertainty from training, generator sampling, corpus construction, or human-AI pairing. They complement the between-run standard deviations in Table 2 and Table 1. All values use the final model for each architecture and follow the reporting convention of Table 3.

For both matched model pairs, the false-positive rate remains higher under the joint procedure on the Qwen3.5-generated texts, although the magnitudes differ.

Paired McNemar tests use predictions on the same DeepSeek-generated AI $\times$ non-sarcastic texts. The discordant-pair counts are 1,021 / 7 for plain dual vs. plain single, 239 / 15 for RCNN dual vs. RCNN single, 805 / 15 for plain dual vs. RCNN dual, and 203 / 100 for the interaction-projection dual-head model vs. the RCNN dual-head model (all $p < .001$). The discordant counts and absolute rate differences provide the substantive effect-size




Figure 2: Sarcasm false-positive rates on AI-generated texts prompted to be non-sarcastic, for five detectors and two generators. Error bars are Wilson 95% intervals conditional on each fixed model.

Model	DeepSeek	Qwen3.5
Plain dual-head	86.3 [84.5, 88.0]	91.3 [89.8, 92.6]
Interaction-projection dual-head	40.9 [38.5, 43.4]	29.5 [27.3, 31.8]
RCNN dual-head	34.1 [31.8, 36.5]	22.1 [20.1, 24.3]
Plain single-head	19.3 [17.4, 21.4]	7.9 [6.6, 9.3]
RCNN single-head	19.3 [17.4, 21.4]	10.6 [9.2, 12.3]

Table 11: Wilson 95% intervals for the sarcasm false-positive rate on AI-generated texts prompted to be non-sarcastic (two generators; %).

information.

In both architectures and both language strata, the false-positive rate is higher under the joint procedure than under its matched sarcasm-only control. Language remains coupled with corpus, genre, and annotation setting, so the stratified results do not isolate a language effect.

The final RCNN dual-head model’s MGTD F1 is 0.938 [0.931, 0.945] on the full set, 0.961 [0.952, 0.969] in Chinese, and 0.923 [0.912, 0.933] in English. Because the human-only subset contains no AI-authored examples, an MGTD F1 score on that subset would not measure source discrimination; Table 7 therefore reports human recall directly.

Chinese sarcasm F1 is higher than English sarcasm F1 for the RCNN model (0.819 vs. 0.429). This contrast also reflects differences in prevalence, genre, and annotation procedure: sarcasm prevalence is 63.8% in Chinese but 14.2% in English, and the English human side is dominated by author-labeled intended-sarcasm tweets from iSarcasmEval whereas the Chinese side uses topic-comment-style ToSarcasm. The positive-class F1 values should therefore not be interpreted as an isolated language effect.

The separately retrained final RCNN model’s

overall sarcasm F1 is 0.663, compared with 0.629 on human-authored texts alone, because the overall score also includes AI-generated texts. This point estimate differs from the five-run mean in Table 1, which aggregates fold-based runs. Within the human subset, absolute joint-versus-sarcasm-only differences in sarcasm F1 are below 0.03 and vary in direction. The negative-transfer effect is therefore concentrated in the quadrants containing AI-generated texts, consistent with §5.3.

False-positive rates use a threshold of 0.5. For the plain dual-head model on AI × non-sarcastic text, the median sarcasm probability is 0.767 with an interquartile range of 0.635–0.839. The distribution therefore lies mostly above the threshold rather than placing only a few cases near it. Single-head medians are 0.204 for the plain model and 0.069 for RCNN.

The median differences above show that joint training raises sarcasm scores on AI-generated texts, while the similar AUC values across models (0.82–0.88) show that their internal ordering is largely preserved. This pattern motivates the representation analyses in Appendix E.

D Illustrative Cases of Disagreement between Detector Probabilities and LLM-Judge Ratings

Illustrative cases were manually selected after inspecting the paired table of RCNN-detector probabilities and ratings from the three LLM judges to show large disagreements; no formal ranking or top- k selection was performed. The selection is non-systematic, and the cases should not be used to estimate the prevalence of any disagreement type. Detector probabilities range from 0 to 1, and LLM-judge ratings are averaged across the three evaluators on a 1–5 scale. Long texts are shortened for display, while Chinese examples retain the original model input and provide Hanyu Pinyin transliterations and English translations in brackets.

The three groups show, respectively, high predicted human-authorship probabilities with low LLM-judge ratings of human-likeness, high predicted sarcasm probabilities with low LLM-judge ratings of sarcasm, and the reverse sarcasm disagreement. In the first human-likeness case, AI-generated text with an obvious reasoning-style preface is nevertheless assigned predicted human-authorship probability 1.00 by the MGTD head, corresponding to the disagreement discussed in §6.

Model	Lang.	n	Dual FP; rate [CI]	Single FP; rate [CI]	Δ pp
Plain	en	1,174	987; 84.1 [81.9, 86.1]	242; 20.6 [18.4, 23.0]	+63.5
Plain	zh	339	319; 94.1 [91.1, 96.1]	50; 14.7 [11.4, 18.9]	+79.4
RCNN	en	1,174	445; 37.9 [35.2, 40.7]	247; 21.0 [18.8, 23.5]	+16.9
RCNN	zh	339	71; 20.9 [17.0, 25.6]	45; 13.3 [10.1, 17.3]	+7.7

Table 12: Language-stratified sarcasm false-positive rates on AI-generated texts prompted to be non-sarcastic (%; Wilson 95% intervals). Dual FP and Single FP give the number of false positives under the dual- and single-head models; Δ is the dual rate minus the single rate. Exact paired McNemar tests give $p < .001$ for both plain comparisons and RCNN English, and $p = 2.4 \times 10^{-5}$ for RCNN Chinese.

Model	Full set	Chinese	English	Human subset	Human $\times$ zh	Human $\times$ en
Plain dual-head	0.556	0.746	0.312	0.586	0.710	0.374
Interaction-projection dual-head	0.643	0.781	0.420	0.593	0.739	0.366
RCNN dual-head	0.663 [.646, .680]	0.819 [.802, .836]	0.429 [.399, .458]	0.629 [.605, .653]	0.774 [.748, .798]	0.399 [.355, .440]
Plain single-head	0.661	0.750	0.514	0.612	0.719	0.434
RCNN single-head	0.664	0.758	0.511	0.603	0.709	0.434

Table 13: Per-language and human-subset sarcasm F1 (final models; bootstrap 95% intervals included for RCNN).

Language	Model condition	Detector probability	Mean LLM-judge rating	Text (excerpt)
Chinese	SFT	1.00	2.67	“好，我需要针对三星回怼华为的新闻生成一条讽刺评论。首先……” [Hao, wo xuyao zhendui Sanxing huidui Huawei de xinwen shengcheng yi tiao fengci pinglun. Shouxian...; “Okay, I need to generate a sarcastic comment on the news of Samsung hitting back at Huawei. First...”] (reasoning-style preface)
Chinese	SFT	1.00	2.33	“跑分竞赛：参数焦虑：跑分焦虑：水货：哈哈哈” [Paofen jingsai: canshu jiaolv: paofen jiaolv: shuihuo: hahaha; “Benchmark race: spec anxiety: benchmark anxiety: grey-market goods: hahaha”] (degenerate output)
English	SFT	1.00	2.67	“Same as me, I’m not sure” (vacuous short sentence)

Table 14: Cases with high predicted human-authorship probability but low LLM-judge ratings of human-likeness. The detector column reports $P(\text{human})$.

Language	Model condition	Detector probability	Mean LLM-judge rating	Text (excerpt)
Chinese	SFT	1.00	2.67	“三星：我才是第一！你看看我这摄像头总分，109+96，112+89，我赢了！” [Sanxing: wo cai shi diyi! Ni kankan wo zhe shexiangtou zongfen, 109+96, 112+89, wo ying le!; “Samsung: I’m the real number one! Look at my camera total scores, 109+96, 112+89, I win!”] (numbers repeated without sarcastic transformation)
Chinese	SFT	0.99	2.67	“恭喜恭喜，美国牛牛牛，继续加油” [Gongxi gongxi, Meiguo niu niu niu, jixu jiaoyou; “Congratulations, congratulations, America is awesome awesome awesome, keep it up”] (formulaic mock-praise)
English	SFT	0.93	2.67	“2 hours left to tip it over, I’m going to the toilet”

Table 15: Cases with high predicted sarcasm probability but low LLM-judge ratings of sarcasm. The detector column reports $P(\text{sarcastic})$.

Language	Model condition	Detector probability	Mean LLM-judge rating	Text (excerpt)
Chinese	Base	0.14	4.00	“特朗普承诺 120 亿补贴像撒大饼。农民收到的却是缩水成硬币的承诺……” [Telangpu chengnuo 120 yi butie xiang sa dabing, nongmin shoudao de que shi suoshui cheng yingbi de chengnuo...; “Trump promised a 12-billion subsidy as if tossing out big pancakes, yet what farmers receive are promises shrunk into coins...”] (metaphorical irony)
Chinese	Base	0.24	4.00	“特斯拉：价格跳水后又跳水。消费者：钱包跳水后又跳水。” [Tesla: jiage tiaoshui hou you tiao shui, xiaofeizhe: qianbao tiaoshui hou you tiaoshui; “Tesla: prices dive and then dive again; consumers: wallets dive and then dive again.”] (parallel-structure irony)
English	Base	0.12	4.00	“Thank you so much for this *beloved* review, Mandy! I’m *overjoyed* you enjoyed it” (asterisk emphasis)

Table 16: Cases with low predicted sarcasm probability but high LLM-judge ratings of sarcasm. The detector column reports $P(\text{sarcastic})$.

Together, these cases motivate candidate hypotheses involving surface markers such as mock-praise formulas, exclamation marks, and colloquial short sentences, as well as semantic reversal, whose prevalence and causal effects require systematic testing. The examples illustrate several disagreement types compatible with the near-zero aggregate Chinese sarcasm correlation in §6.

E Representation-Level Analysis of Error Distribution between Source Groups

This appendix provides additional protocol details, diagnostic results, and implementation checks for the representation-level tests in §5.3.2.

E.1 Setup. All analyses use the frozen plain dual-head model that produces the 86.3% result reported in the main text. We divide the samples into two halves: one half is used to estimate the direction, and the other is used to compute metrics. The appendix baseline is therefore 84.8% on the evaluation half rather than 86.3% on the full set.

Implementation checks show that manually computed logits match model outputs with a maximum error of 1.4×10^{-6} and that the residual after the erasure projection is about 10^{-6} . LEACE satisfies its defining property: after erasure, the dif-

ference between class means on the fitting half has a relative residual of 7×10^{-6} , and a linear probe evaluated on that half is at chance (AUC 0.514). This in-sample result serves only as an implementation check, whereas Table 17 reports probe performance after erasure on the held-out half. The AUC of the source probe is 0.986 with true labels, and a permuted-label sanity check gives 0.512. These probe values refer to different evaluation sets: 0.986 is the five-fold result on the unmodified full representation set, 0.982 is the corresponding baseline on the held-out half used to evaluate the interventions, 0.514 is the in-sample LEACE implementation check, and 0.589 is probe performance after erasure on the held-out half.

E.2 Protocol. The linear probe is a logistic-regression classifier evaluated with five-fold cross-validation. As a sanity check, a permuted-label control retrain the same probe after shuffling source labels, and correlation analyses are repeated within text-length tertiles. For alignment, Σ is the Ledoit–Wolf shrinkage covariance matrix (Ledoit and Wolf, 2004) fitted to all pooled DualTest representations, with eigenvalues clipped below 10^{-8} ; each w is a task-head weight difference or probe direction. The whitening transform is transductive: its covariance is estimated from unlabeled representations in both halves. For two such directions w_1 and w_2 , whitened alignment is $\cos(\Sigma^{1/2}w_1, \Sigma^{1/2}w_2)$. We report absolute cosine because the sign of a probe direction depends on the source-label convention.

Erasure tests cover one direction, INLP subspaces with $k = 1$ –16 removed directions (Ravfogel et al., 2020), and LEACE (Belrose et al., 2023). Controls use ten sampled random whitened directions and directions estimated from permuted labels. We construct the source-direction-orthogonal first principal component (srcORTH-PC1) by projecting the representations in the fitting half onto the orthogonal complement of the estimated linear source direction and then extracting the maximum-variance direction in that projected space. Erasure removes each representation’s scalar projection onto this direction. This operation leaves the measured linear source decodability unchanged, although orthogonality to one estimated direction does not imply independence from all multidimensional or nonlinear source information.

E.3 Diagnostic evidence. Source is highly linearly decodable: AUC is 0.986 overall, 0.990 on sarcastic texts, and 0.985 on non-sarcastic texts. The permuted-label control gives 0.512. The single-head model also yields an AUC of 0.974, showing that source decodability is a background property shared by both models. On AI $\times$ non-sarcastic texts, the MGTD-head AI logit and sarcasm-head logit correlate at $\rho = 0.519$; correlations by text-length tertile are 0.44, 0.56, and 0.63.

The sarcasm and source directions are nearly orthogonal in the original geometry ($\cos \approx -0.03$). In the whitened representation space, the absolute cosine is 0.41 for the probe-estimated source direction and 0.49 for the MGTD-head source direction. These correlations and directional alignments link source-related structure to the affected sarcasm decisions and motivate the intervention tests below.

E.4 Intervention evidence. We define gap as the AI $\times$ non-sarcastic false-positive rate minus the human $\times$ non-sarcastic false-positive rate. Source AUC measures the performance of a probe re-trained after erasure.

Erasure scheme	AI $\times$ n FPR	H $\times$ n FPR	Gap	AUC
None (baseline)	84.8	32.8	52.0	0.982
Single dir. (head / probe / random)	84.8–86.2	—	—	≈ 0.98
INLP ($k = 16$)	89.8	—	—	0.974
LEACE (linear)	68.7	49.5	19.1	0.589
Random	84.9	32.7	—	0.982
whitened $\times 10$				
srcORTH-PC1	66.7	54.1	12.6	0.982

Table 17: Selected intervention summary. H $\times$ n = human $\times$ non-sarcastic; AI $\times$ n = AI $\times$ non-sarcastic; FPR in %; AUC is the AUC of the source probe after erasure. Bold marks the two diagnostic interventions discussed in §5.3.2, not best performance.

The summarized single-direction and INLP rows do not show a reduction in the AI-generated subgroup comparable to LEACE. LEACE reduces measured linear source decodability and the gap defined above from 52.0 to 19.1. Its gap change is not matched by any of the ten sampled random whitened directions.

The srcORTH-PC1 intervention shows that reducing measured linear source decodability is not necessary for changing how errors are distributed between source groups. It leaves measured source decodability intact but changes the AI-generated subgroup rate to 66.7 and the difference between source groups to 12.6, an effect similar to LEACE

and unmatched by any of the ten sampled random directions. The intervention identifies a direction whose semantic content and role in the transfer observed during training remain targets for further causal analysis.

We use an unweighted group mean to give the two non-sarcastic source groups equal importance despite their different sizes. Under LEACE, this mean false-positive rate is almost unchanged (58.8 $\rightarrow$ 59.1). This result shows that the affected directions shape how non-sarcastic errors are allocated between AI and human texts.